



LUDWIG-
MAXIMILIANS-
UNIVERSITÄT
MÜNCHEN

Deep Learning for Physicists

Lecture #1: Introduction

Kosmas Kepesidis

Outline

- About the course
- Introduction to artificial intelligence (AI)
- Models for data analysis
- Building blocks of neural networks

About the course

About the course

- Things to keep in mind:
 - Comprehensive introductory course on deep learning – adapted for physicists
 - Target audience: BSc & MSc students of physics
 - Prior knowledge on the topic is not required!
 - Prerequisites: basic math and physics from the BSc in physics and some programming experience (preferably with Python)

About the course

- Course content:
 1. Basics of artificial neural networks
 2. Architectures of deep neural networks
 3. Interpretability, uncertainties, objectives

About the course

- Course content:

- Basics

- Data and analysis in physics
 - Building blocks of neural networks
 - Optimization of neural networks (model training)
 - Practical methodology (basics of machine-learning engineering)

About the course

- Course content:
 - Architectures of deep neural networks
 - Fully-connected neural networks
 - Convolutional neural networks (CNNs)
 - Recurrent neural networks (RNNs)
 - Graph networks
 - Hybrid architectures

About the course

- Course content:
 - Interpretability, uncertainties, objectives
 - Interpretability
 - Uncertainties
 - Objective functions: deeper insights

About the course

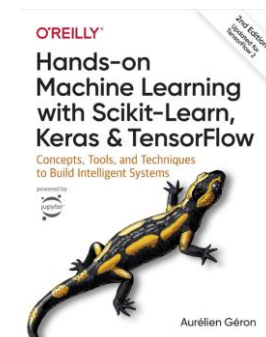
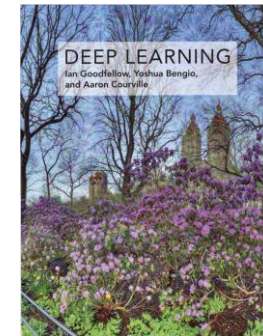
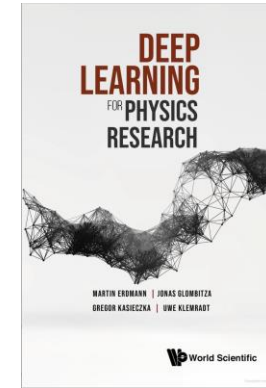
- Next semester: Advanced Deep Learning for Physicists

- Advanced topics

- Unsupervised and semi-supervised learning
- Autoencoders
- Generative models
- Outlier and anomaly detection
- Special topics
- Mathematical foundations

About the course

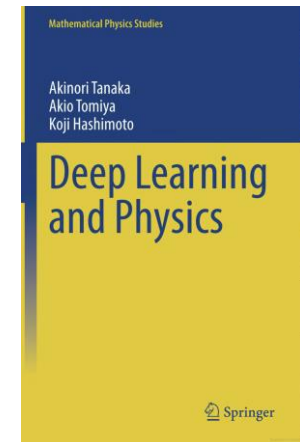
- Textbooks used for the lectures:
 - **Deep Learning for Physics Research**, Gregor Kasieczka, Jonas Glombitza, and Martin Erdmann (2021)
 - **Deep Learning**, Ian Goodfellow, Yoshua Bengio, Aaron Courville (2016)
 - Open access: <https://www.deeplearningbook.org/>
 - **Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow** (2nd Edition), Aurélien Géron (2019)



About the course

- Another textbook... not used here but still interesting!

➤ **Deep Learning and Physics**, Akinori Tanaka, Akio Tomiya, Koji Hashimoto (2021)



About the course

- Tutorials:
 - Problems sets
 - Computational assignment – Jupyter notebook in Python
 - Assignments will not contribute to the grade
- Written final exam :
 - Final exam will contribute to 100% of your grade

About the course

- Course material:

- Moodle: “Deep Learning for Physicists” <https://moodle.lmu.de/course/view.php?id=25091>

- Lecture slides
 - Problem sets (+ solutions)
 - Jupyter notebooks (+ solutions)
 - Useful papers

About the course

- Software:

- Python language: <https://docs.python.org/3.7/tutorial/index.html>
 - Would be useful to get familiar within the first two weeks of the course
- Anaconda distribution (open source): <https://www.anaconda.com/>
 - Recommend to download it – will be used in the tutorials



Introduction to artificial intelligence (AI)

Introduction to artificial intelligence (AI)

- The field of AI is driven by the desire to create machines that think
- In the early days of AI, problems that are difficult for humans and easy for computers were tackled – **rule-based systems** (RBS)



- Early success based on RBS: IBM's Deep Blue chess computer defeated world champion Garry Kasparov in 1997 (Hsu, 2002)

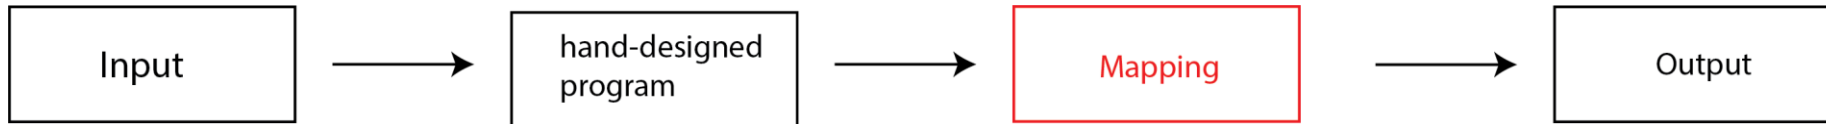
Introduction to artificial intelligence (AI)

- The real challenge of AI remained to tackle problems that are straightforward to humans
- **Knowledge based** approach to AI: knowledge about the world hard-coded in a computer program
- Cyc (Lenat and Guha, 1989): inference engine + database of statements
- Memorizes large collections of general knowledge
- Has proven challenging in practice (Linde, 1992)



Introduction to artificial intelligence (AI)

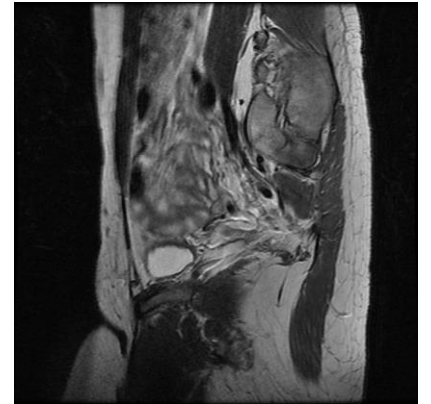
- Solution: ability to acquire own knowledge by discovering patterns in data
- This approach is known as **machine learning** (ML)
- Examples of *classical* ML algorithms: logistic regression, decision trees, support vector machines, naïve Bayes



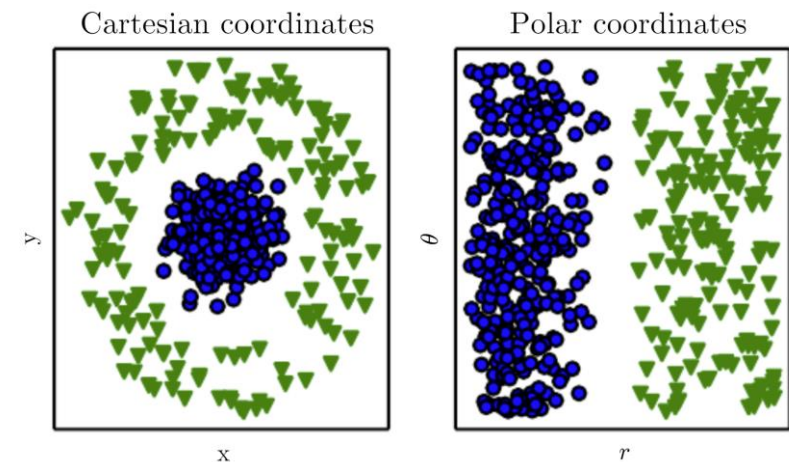
- Examples of problems ML solves:
 - recommend cesarean delivery (Mor-Yosef et al., 1990)
 - separate legitimate e-mail from spam e-mail

Introduction to artificial intelligence (AI)

- ML works well with predefined/preselected **features** (or **observables**) selected by researcher
- Limitation of ML: no influence on feature definition
 - Example: ML-based cesarean-delivery recommender depends on the doctors report (features), doctor needs to interpret MRI scan of patient

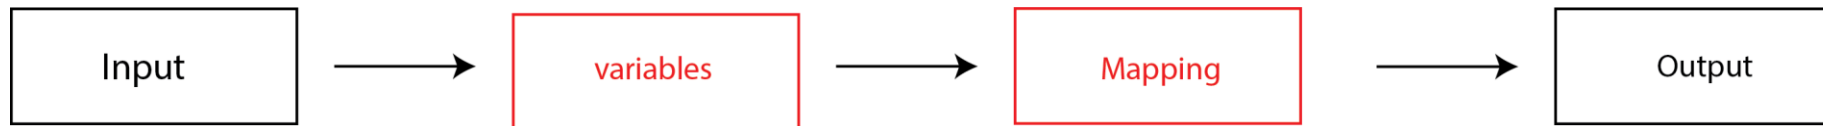


- **Representations** play crucial role in ML



Introduction to artificial intelligence (AI)

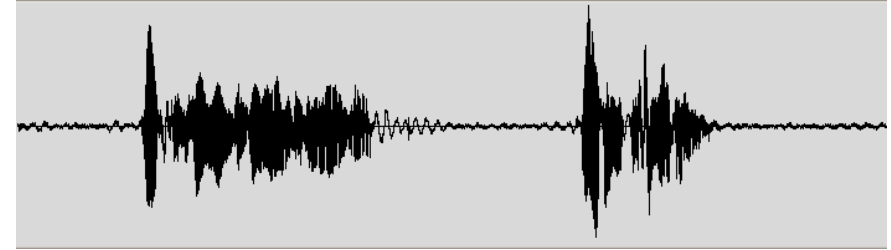
- Use machine learning to discover not only the mapping from representation but also the representation itself
- Approach known as **representation learning**



- Example from physics: extracting velocity of particle from location and time data
- Representation learning can be challenging – need to separate the **factors of variation** (FOV), which can be of high abstraction

Introduction to artificial intelligence (AI)

- Example: speech recording
 - FOV: age, gender, accent



Source: <https://cmusphinx.github.io/>

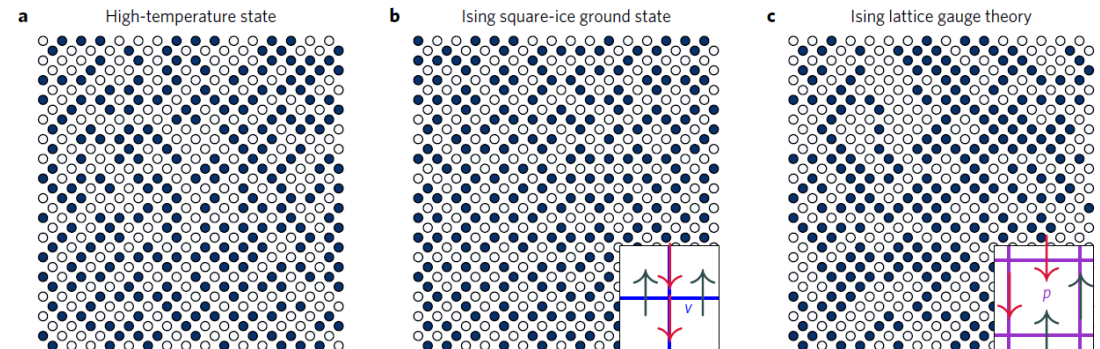
- Example: image of a car
 - FOV: position of the car, its color, angle and brightness of the sun



Source: <https://en.wiktionary.org/>

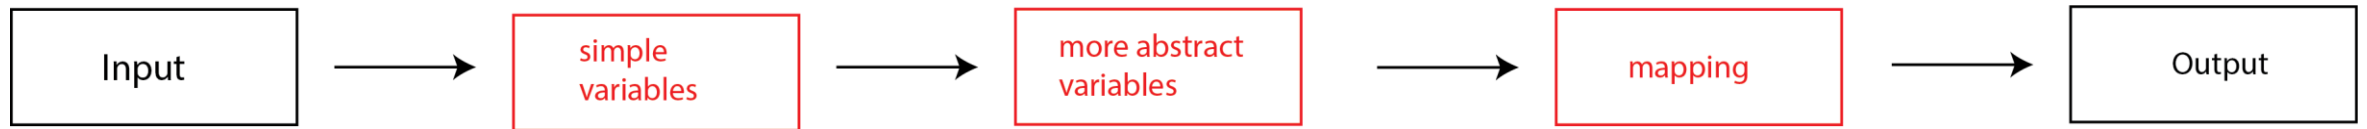
- Example: phases of matter
 - FOV: temperature

Nature Phys **13**, 431–434 (2017)



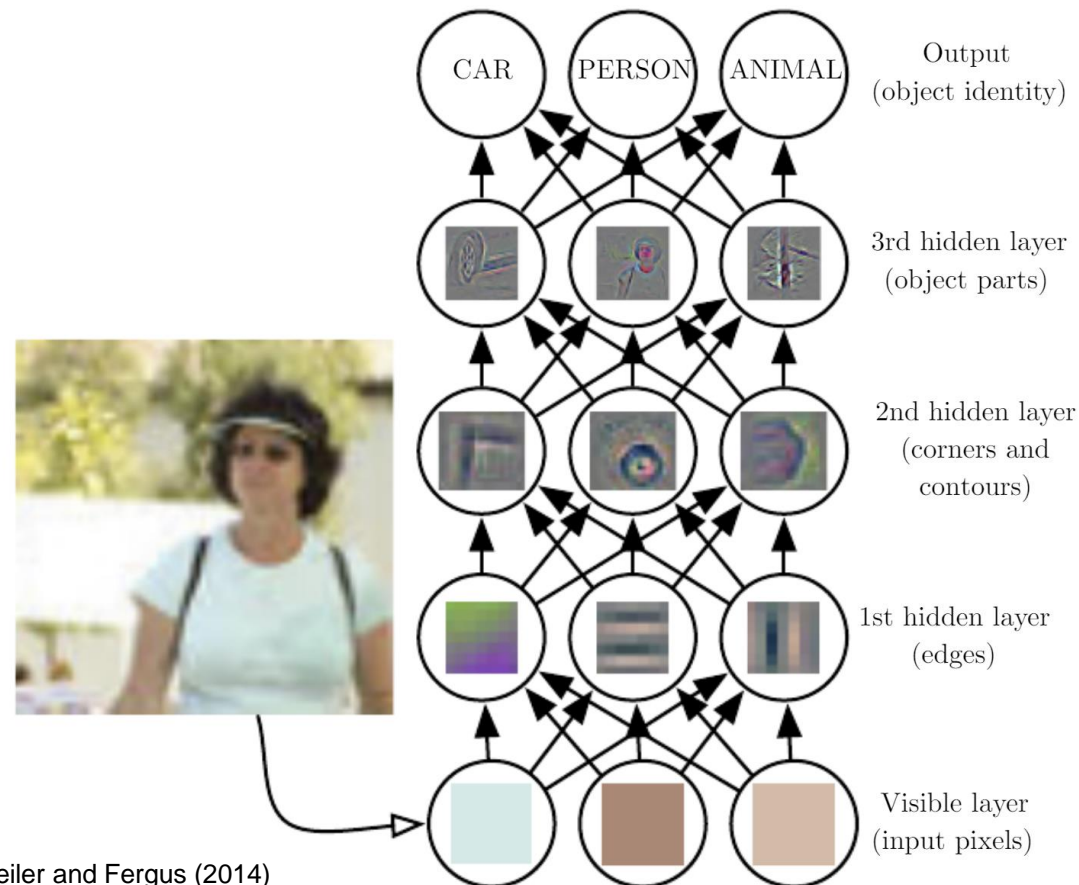
Introduction to artificial intelligence (AI)

- **Deep learning** (DL): constructing more abstract representations from simpler representations *hierarchically*



Introduction to artificial intelligence (AI)

- **Deep learning (DL):** constructing more abstract representations from simpler representations *hierarchically*



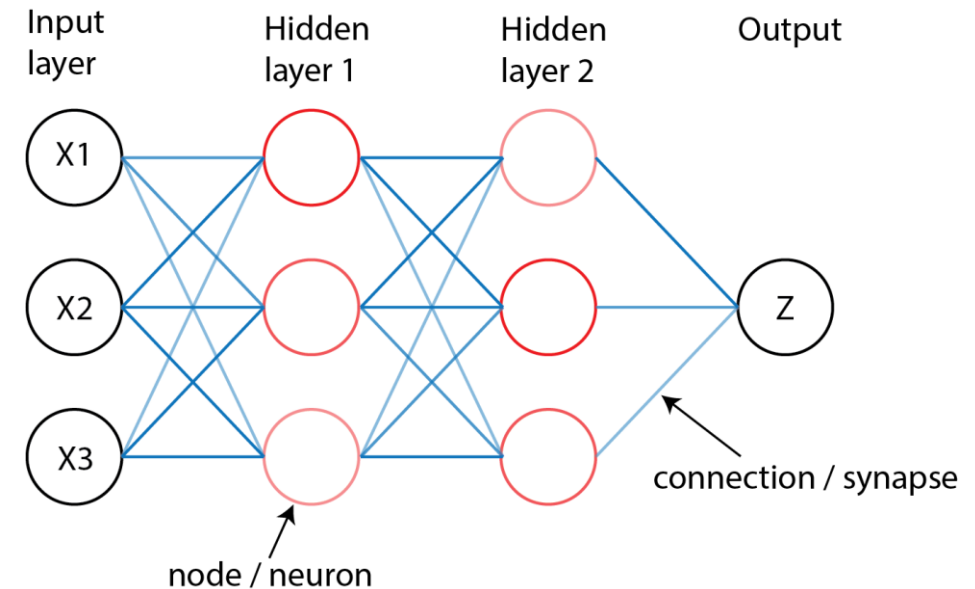
Source: Zeiler and Fergus (2014)

Introduction to artificial intelligence (AI)

- **Artificial neural networks (ANN):**
 - Signal flows only in one direction: feedforward neural networks (FNN)
- **Deep neural networks (DNN):** ANNs with many hidden layers

“It’s deep if it has more than one stage of non-linear feature transformation.”

– Yann LeCun, DL researcher

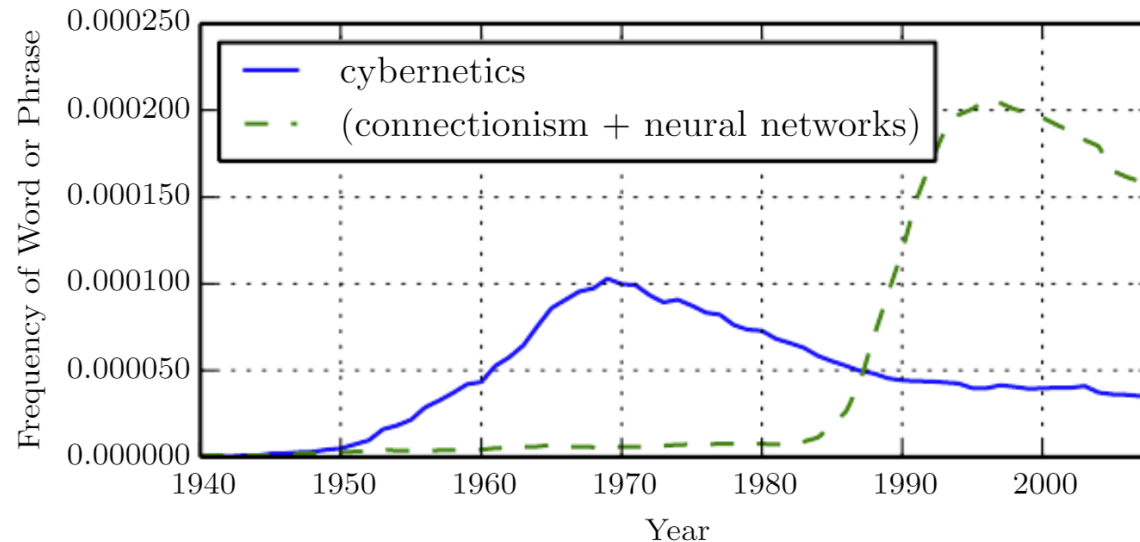


Introduction to artificial intelligence (AI)

- **DNNs** are versatile, powerful, and scalable – can tackle highly-complex tasks such as:
 - Classification of billions of images
 - Speech-recognition services (e.g. Apple's Siri, Amazon's Alexa, etc)
 - Recommend products or videos to millions of users
 - Beat world champions at the game of Go (DeepMind's AlphaGo)
 - Development of self-driving cars
- ...

Introduction to artificial intelligence (AI)

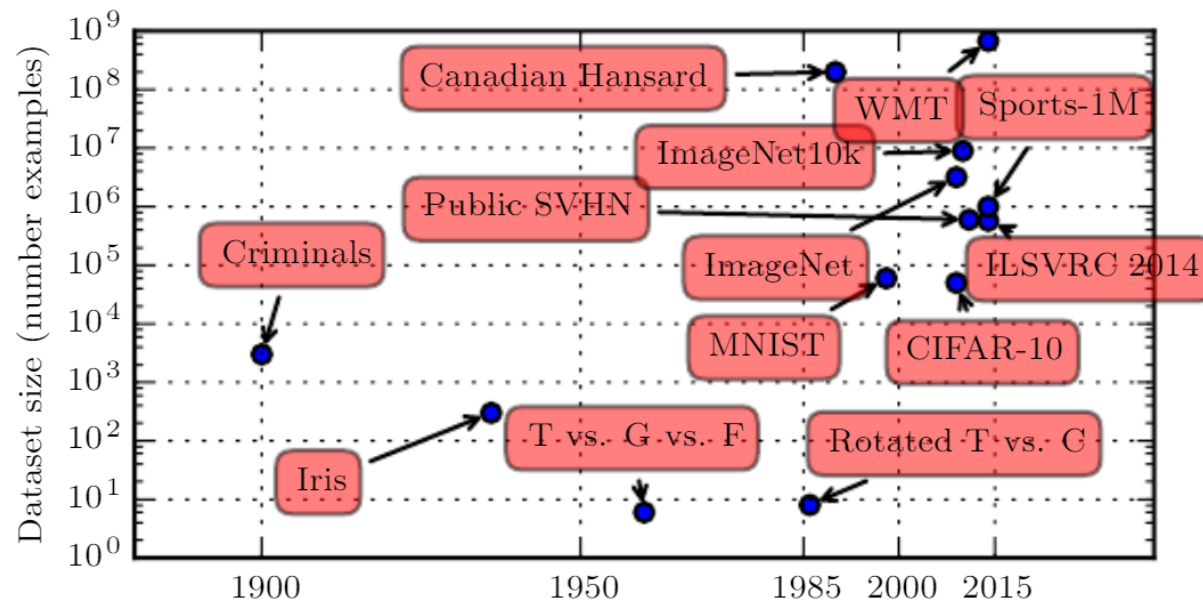
- DL has gone with different names in different time periods:
 - 1940s–1960s: **cybernetics** – theories of biological learning
 - 1980s–1990s: **connectionism** – backpropagation (Rumelhart et al., 1986)
 - 2006-now: **deep learning** – advancements in technology



Source: Goodfellow et al., Deep Learning

Introduction to artificial intelligence (AI)

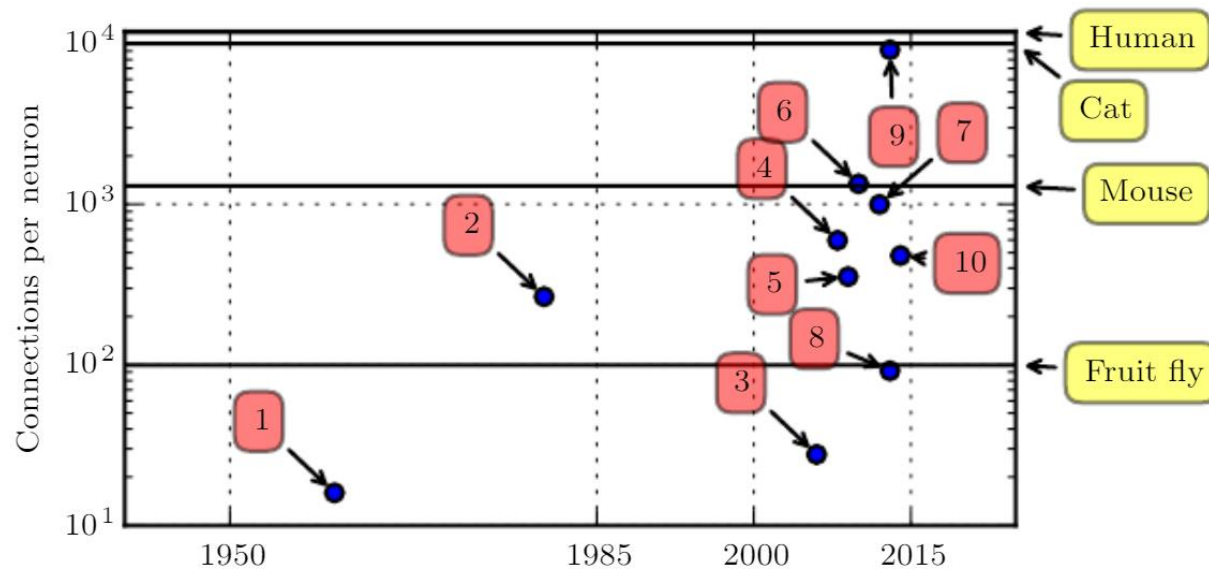
- Causes for rapid developments
 - amount of available training data has increased



Introduction to artificial intelligence (AI)

- Causes for rapid developments

➤ Larger and more complex models developed, tailored to application



1. Adaptive linear element (Widrow and Hoff, 1960)
2. Neocognitron (Fukushima, 1980)
3. GPU-accelerated convolutional network (Chellapilla et al., 2006)
4. Deep Boltzmann machine (Salakhutdinov and Hinton, 2009a)
5. Unsupervised convolutional network (Jarrett et al., 2009)
6. GPU-accelerated multilayer perceptron (Ciresan et al., 2010)
7. Distributed autoencoder (Le et al., 2012)
8. Multi-GPU convolutional network (Krizhevsky et al., 2012)
9. COTS HPC unsupervised convolutional network (Coates et al., 2013)
10. GoogLeNet (Szegedy et al., 2014a)

Introduction to artificial intelligence (AI)

- Causes for rapid developments
 - Hardware advancements such more memory/processing power and access to **Graphics Processing Units** (GPUs) – possibility to train more complex and accurate models

Introduction to artificial intelligence (AI)

- Causes for rapid developments

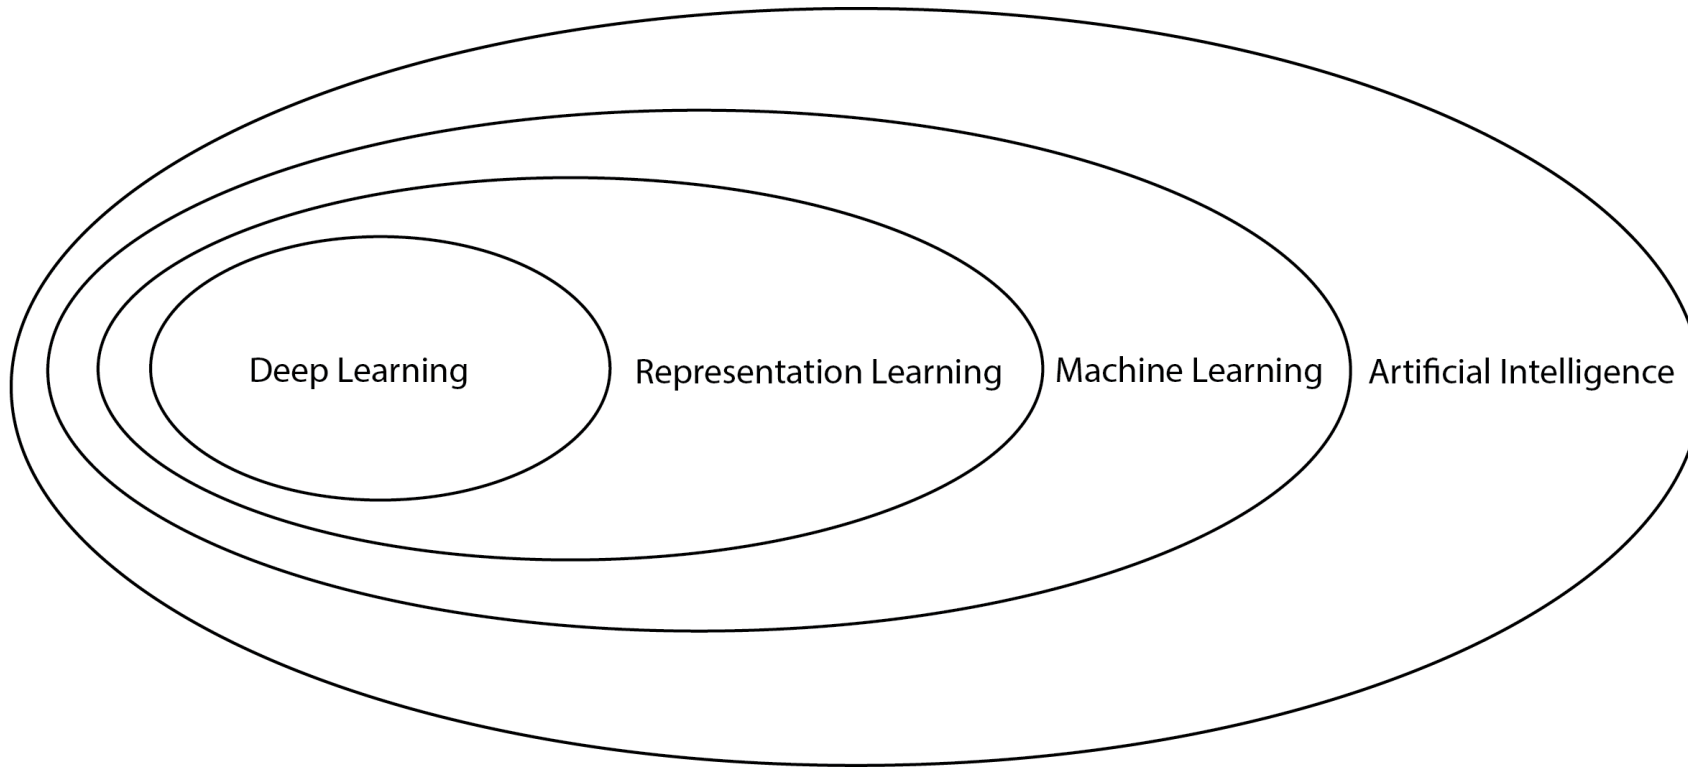
- User-friendly open-source library software for DL:

- Theano
- TensorFlow
- Keras
- PyTorch

```
from tensorflow import keras
layers = keras.layers

model = keras.models.Sequential()
model.add(layers.Dense(4, activation='relu', input_dim=2))
model.add(layers.Dense(1, activation='sigmoid'))
model.compile(loss='MSE', optimizer='SGD', metrics=['accuracy'])
model.fit(xdata, ydata, epochs=200)
```

Introduction to artificial intelligence (AI)



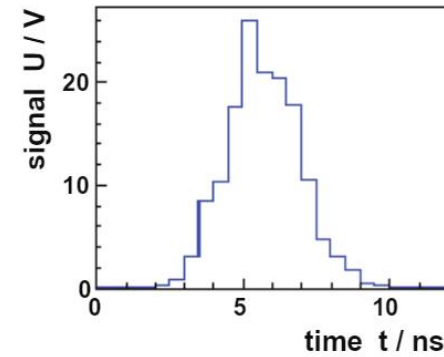
*"In AI, system should be understood as including the human engineers.
Most of the data-generalization conversion happens during model design."*

– Francois Chollet, AI researcher

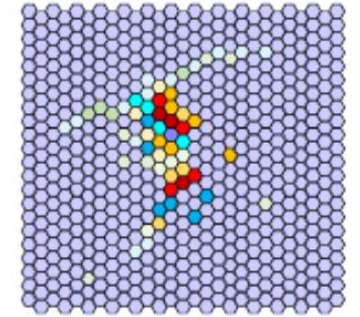
Models for data analysis

Models for data analysis

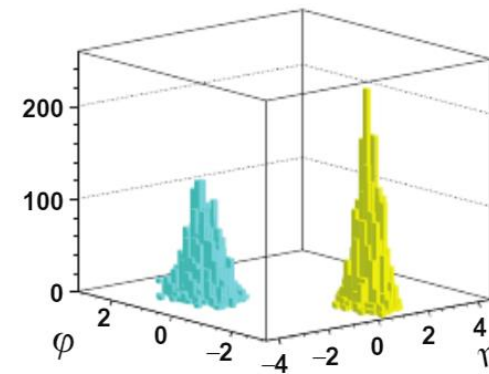
- Examples of data in physics
 - (a) amplitude as a function of time
 - (b) camera image of a telescope
 - (c) particles of a collision
 - (d) sequences of sort pulses



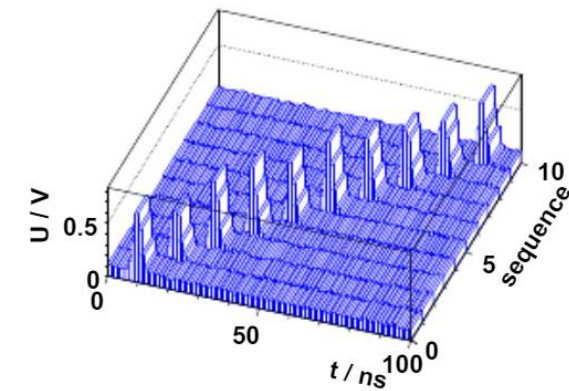
(a)



(b)



(c)



(d)

Models for data analysis

- Data types in physics
 - Data from k events
 - 1D sequences
 - Image-like (grid data)
 - Composed: image-like + 1D sequences
 - Point cloud data (3D, spacetime)
 - Heterogeneous data

Models for data analysis

- Data types in physics

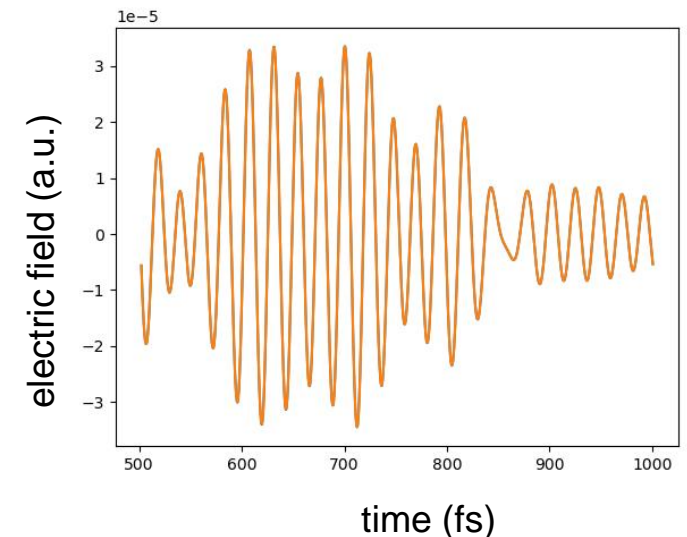
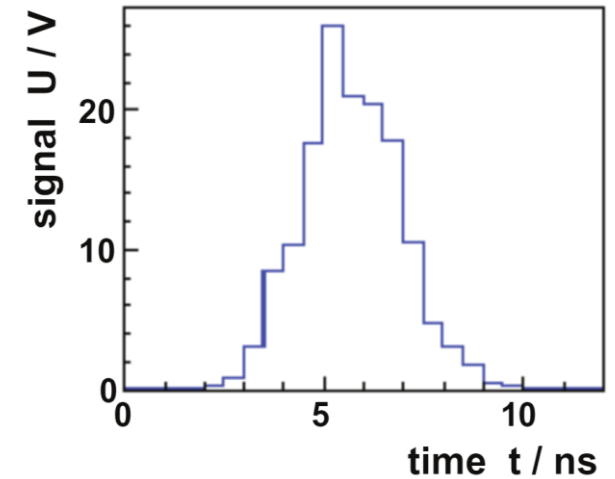
- Data from k events

- General situation of measuring k times n variables
 - Each event corresponds to an n -dimensional vector: \vec{x}
 - This results into an $k \times n$ table (tabular data)

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
0	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8	6
1	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5	6
2	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1	6
3	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	6
4	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9	6

Models for data analysis

- Data types in physics
 - 1D sequences (time series / time traces)
 - Variable x corresponds to an ordered list of numerical values
 - Typically, signal functions $x(t)$
 - But it can be extended to n -dimensional vector $\vec{x}(t)$
 - Examples:
 - a. Time-ordered series of electric amplitudes of a microphone or loud-speaker for recording or reproducing audio
 - b. Sensors in mobile phones such as a pressure sensor, light sensor, temperature sensor, and three-dimensional acceleration sensors
 - c. Electric-field response measured after a laser pulse passes through a liquid medium

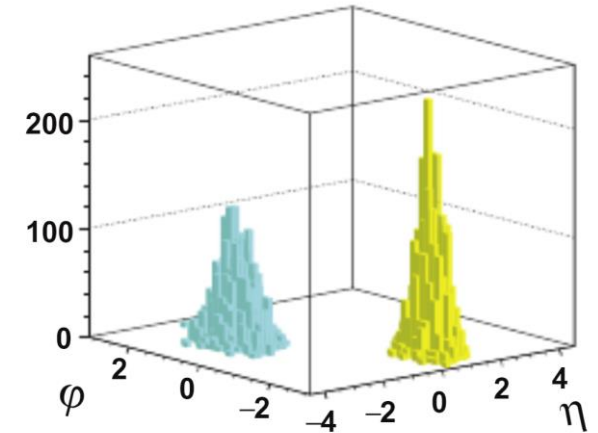
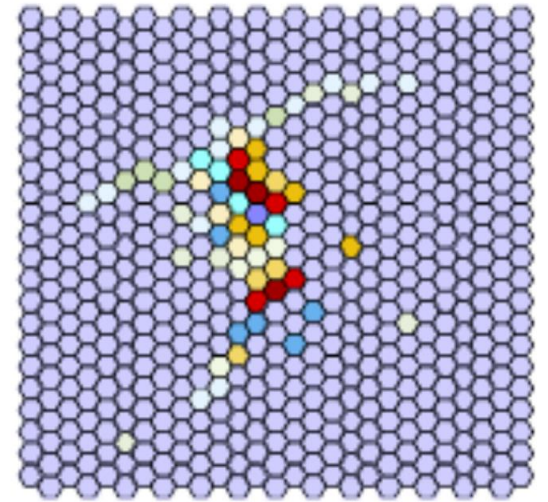


Models for data analysis

- Data types in physics

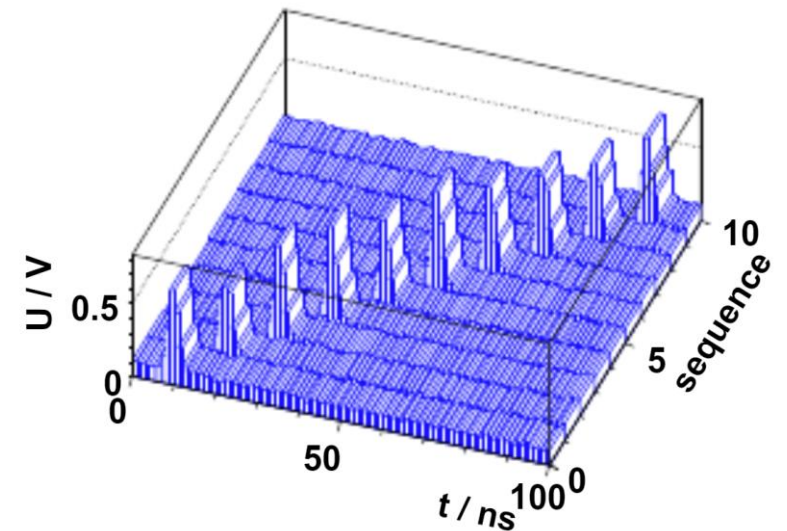
- Image-like (grid data)

- Data on a regular 2D (or higher dimensional) grid
 - Examples include regular digital images: pixels arranged on a 2D Cartesian grid
 - In astronomy such data result from instruments such as telescopic cameras
 - Sensors typically arranged in hexagonal structure instead of Cartesian
 - In particle physics: particles of a collision projected onto a physics-motivated coordinate system



Models for data analysis

- Data types in physics
 - Composed: image-like + 1D sequences
 - Sequences of image-like data
 - Examples:
 - videos, as sequences of images
 - sequences of short pulses

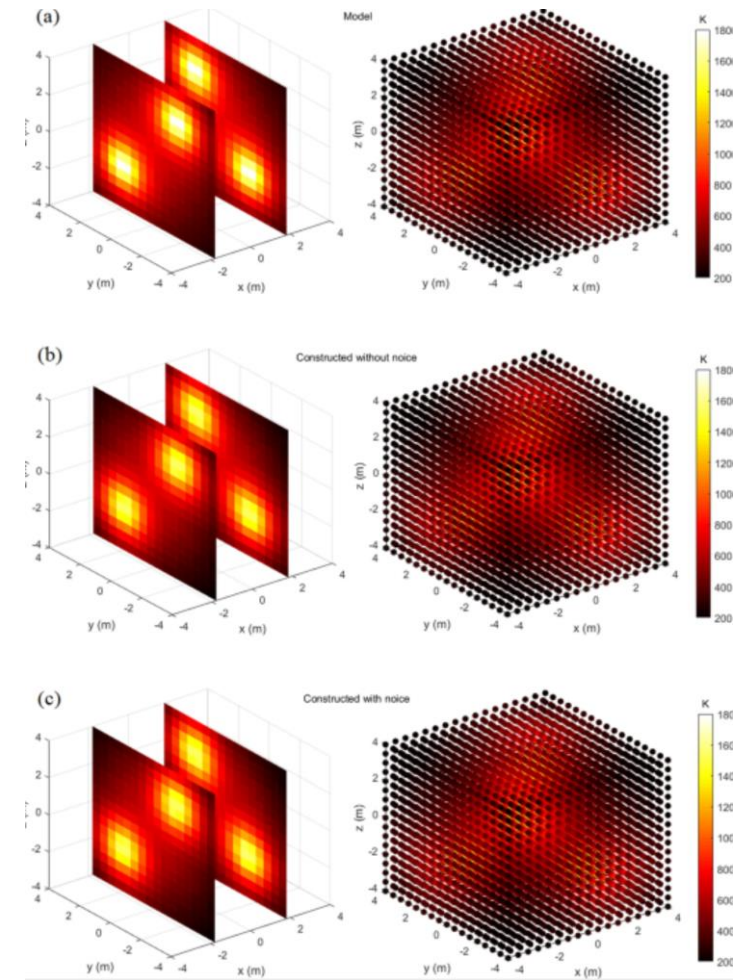


Models for data analysis

- Data types in physics

- Point cloud data (3D, spacetime)

- 3D form/object described as a set of data points in space
 - Each point is related to a vector for locating its position
 - The collection of such vectors form the **point cloud**
 - Examples:
 - a. Temperature measurements in 3D space
 - b. Data for 3D scanners



Source: X. Shen, Construction of three-dimensional temperature distribution using a network of ultrasonic transducers, Scientific Reports 9, 12726 (2019)

Models for data analysis

- Data types in physics

- Heterogeneous data

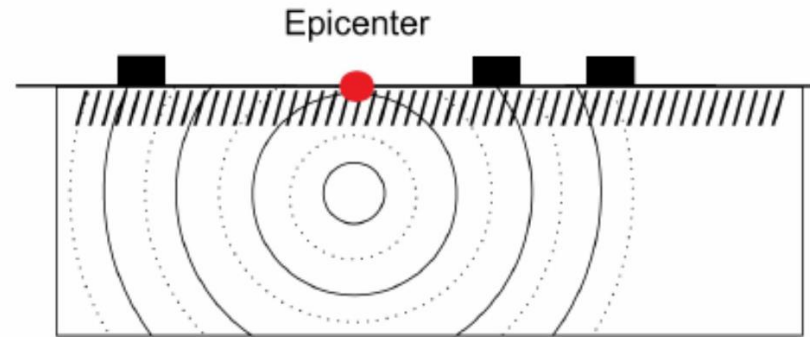
- Sometimes, a measurement is the result of information captured by an *ensemble of devices*
 - Examples:
 - a. In seismology, distributed sensors are used for the detection of earthquake epicenter
 - b. In particle physics, detectors such as ATLAS and CMS at CERN consist of millions of sensors, which comprise different technologies and all together form a highly heterogeneous measurement system

Models for data analysis

- Model building
 - Scientific question themselves determine which type of information and from how many sensors should be used
 - Intuitive understanding of physics, generates expectations about the aspects of measurements that can be utilized for solving a task

Models for data analysis

- Model building: locating earthquake epicenter



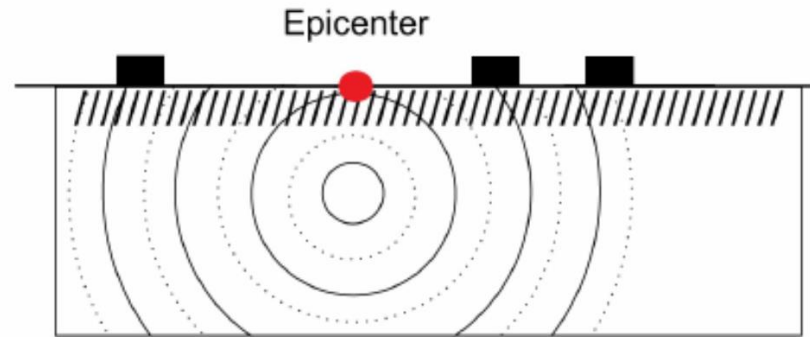
- Physicist's approach:
 - Sheer and compression waves – known velocities
 - Measuring signal times at each station
 - Determine location of epicenters

$$t = x / v_s$$

$$t = x / v_p$$

Models for data analysis

- Model building: locating earthquake epicenter

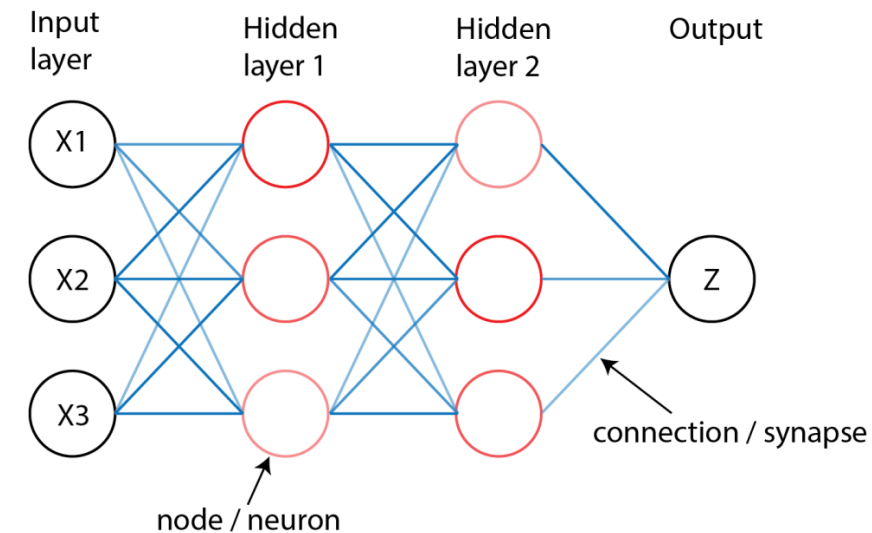


- Data-driven approach:
 - Network becomes the model – relationship is learned from training data
 - No need to specify positions of stations or velocities
 - Input: arrival times $t_{i,p}$ $t_{i,s}$
 - Output: longitude, latitude of epicenter (l,b)
 - Training: iterative process of comparing with the true values of (l,b)

Models for data analysis

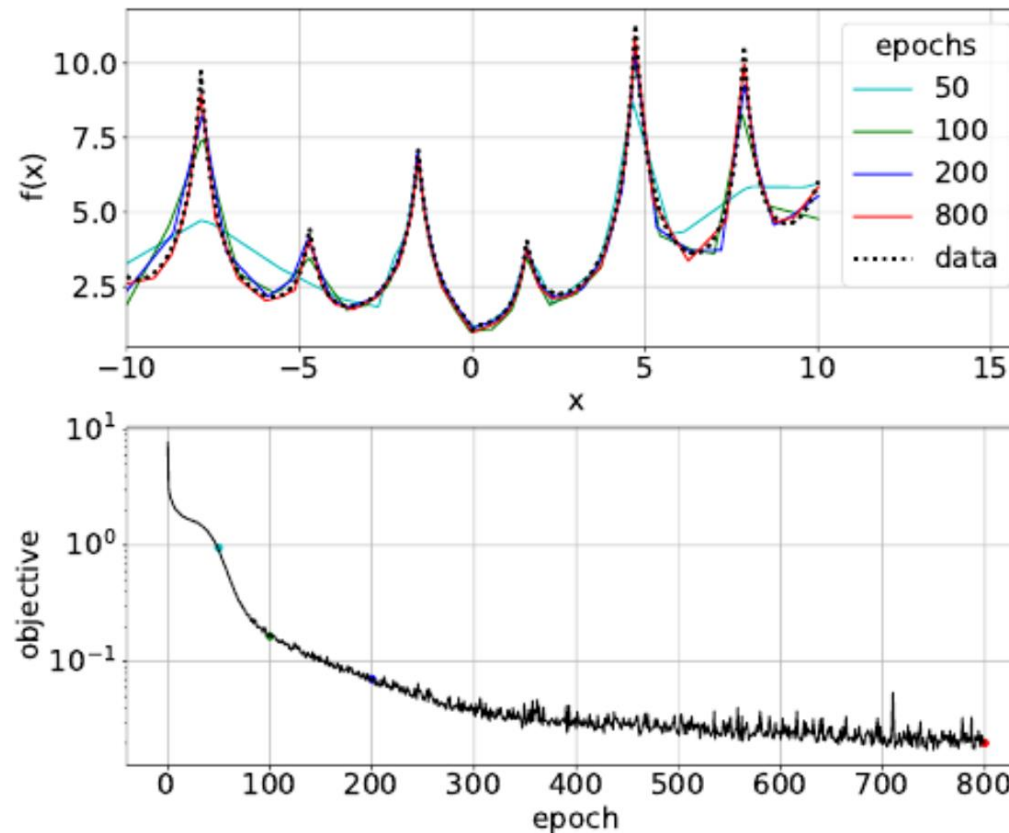
- Data-driven model optimization (network training)

- Network parameters: weights of connections
- Training data set:
 - input values, vector \vec{x}
 - Target values y , labels
- Objective function \mathcal{L} :
 - distance between network predictions $f(\vec{x})$ and target y
 - \mathcal{L} depends on scientific question
 - Common way: mean value of squared residuals $\mathcal{L} = \left(\frac{1}{k}\right) \sum_{i=1}^k \Delta^2$
 - with residuals $\Delta = |f(\vec{x}) - y|$
 - i runs over all data points



Models for data analysis

- Data-driven model optimization (network training)



Building blocks of neural networks

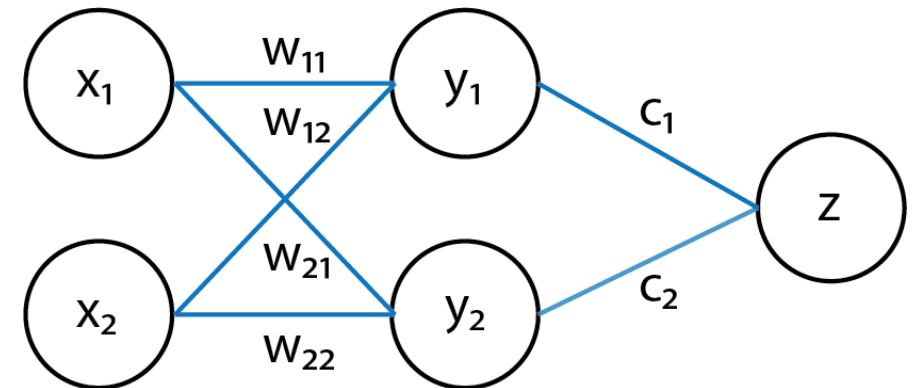
Building blocks of neural networks

- Linear mapping and displacement

$$\vec{y} = \mathbf{W}\vec{x} + \vec{b}$$

$$\begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} \\ W_{21} & W_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

- Affine transformation
 - *Matrix* \mathbf{W} : weights
 - Vector b : bias



Building blocks of neural networks

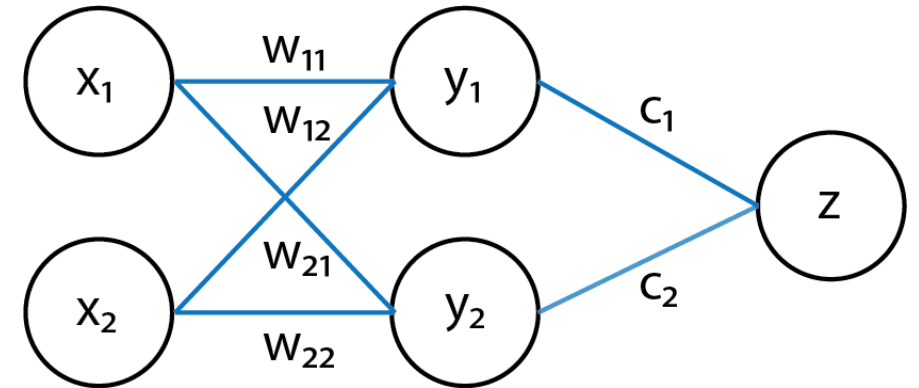
- Linear mapping and displacement

$$\vec{y} = \mathbf{W}\vec{x} + \vec{b}$$

$$\begin{pmatrix} y_1 \\ y_2 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} \\ W_{21} & W_{22} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \end{pmatrix}$$

- Another affine transformation

$$\begin{aligned} z &= \mathbf{W}'\vec{y} + d = (c_1 \quad c_2) \begin{pmatrix} y_1 \\ y_2 \end{pmatrix} + d \\ &= W_1''x_1 + W_2''x_2 + b'' \end{aligned}$$



$$W_1'' = c_1W_{11} + c_2W_{21}$$

$$W_2'' = c_1W_{12} + c_2W_{22}$$

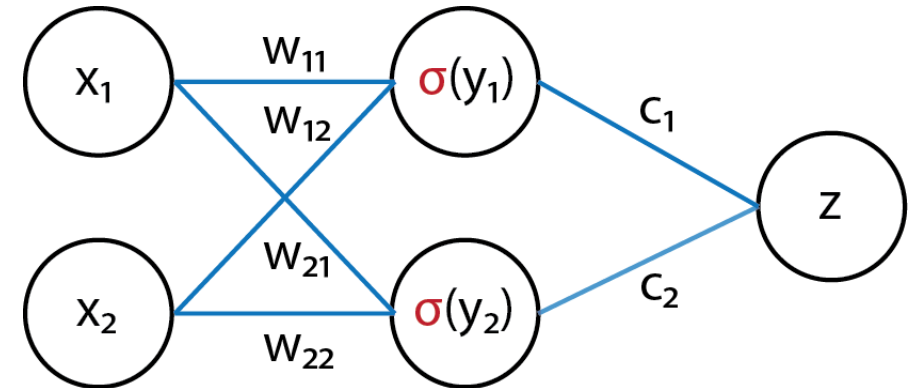
$$b'' = c_1b_1 + c_2b_2 + d$$

Building blocks of neural networks

- Nonlinear mapping: **activation function**

- Nonlinearity in the neurons

$$\sigma(\vec{y}) = \begin{pmatrix} \sigma(y_1) \\ \sigma(y_2) \end{pmatrix}$$



$$z = \mathbf{W}' \sigma(\vec{y}) + d = (c_1 \quad c_2) \begin{pmatrix} \sigma(y_1) \\ \sigma(y_2) \end{pmatrix} + d$$

$$= c_1 \sigma(W_{11}x_1 + W_{12}x_2 + b_1) + c_2 \sigma(W_{21}x_1 + W_{22}x_2 + b_2) + d$$

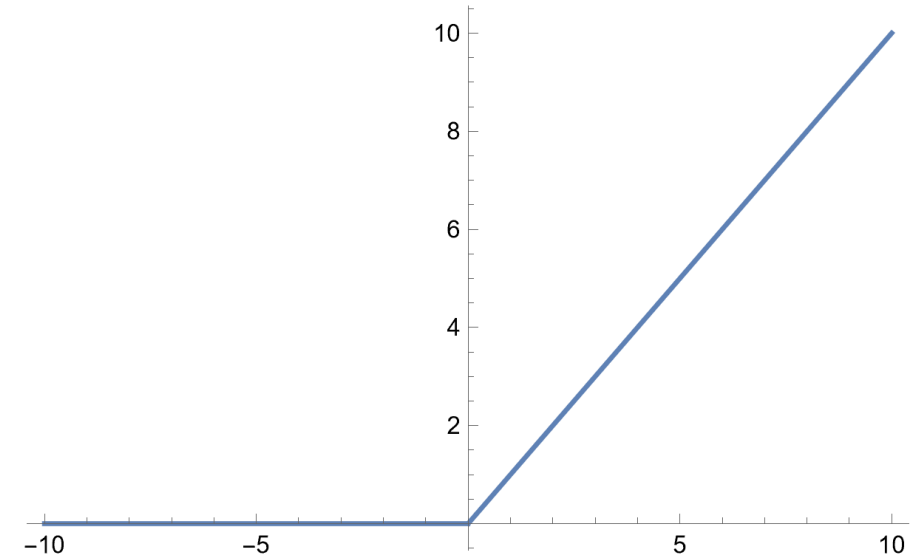
Building blocks of neural networks

- Nonlinear mapping: **activation function**

- Types of activation functions

- Rectified Linear Unit (ReLU)

$$\sigma(x) = \max(0, x) = \begin{cases} x & x \geq 0 \\ 0 & x < 0 \end{cases}$$



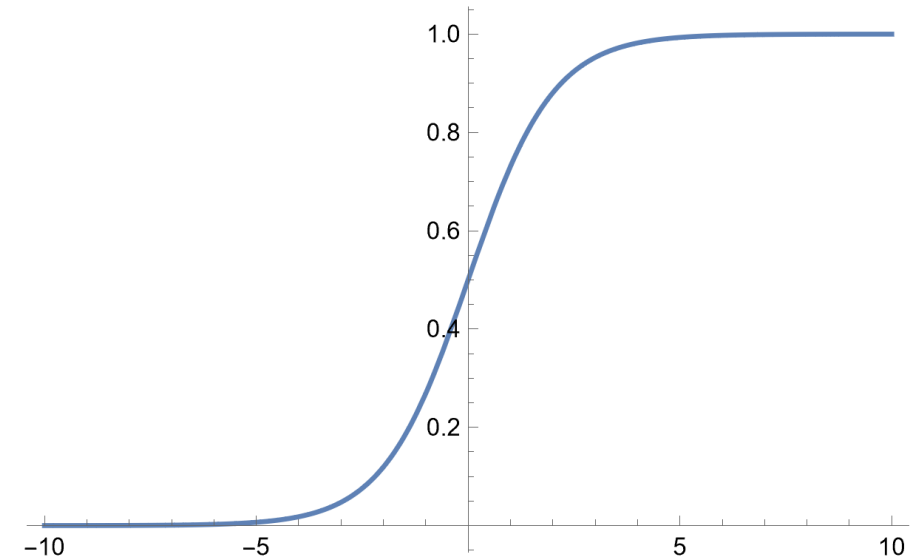
Building blocks of neural networks

- Nonlinear mapping: **activation function**

- Types of activation functions

- Sigmoid (logistic) function

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



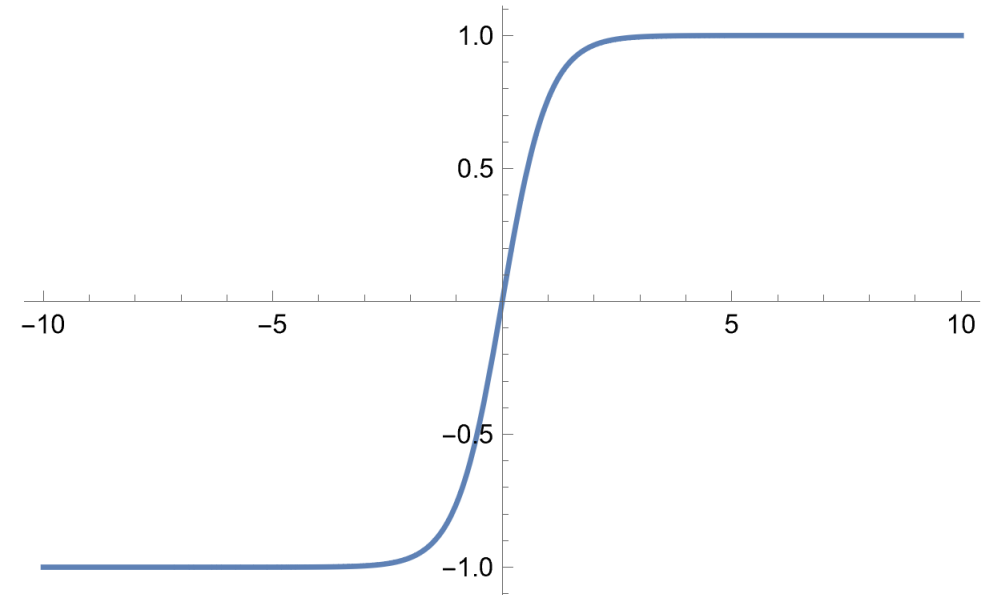
Building blocks of neural networks

- Nonlinear mapping: **activation function**

- Types of activation functions

- Hyperbolic tangent

$$\sigma(x) = \tanh(x)$$



Building blocks of neural networks

- Network predictions

- At each node, two transformations

1. Affine mapping:

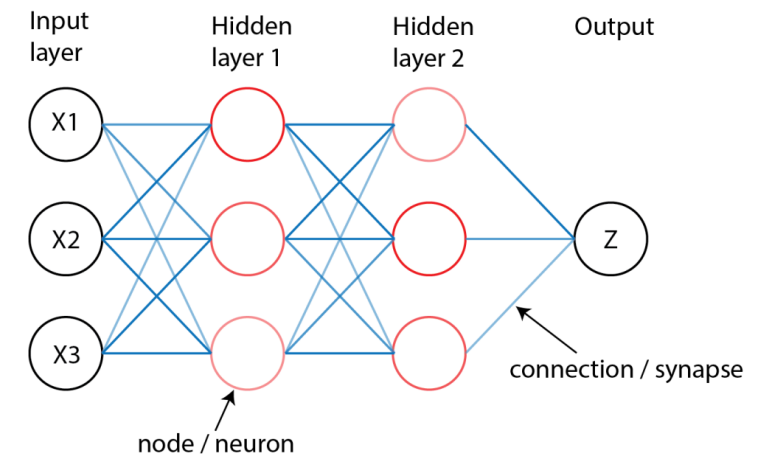
$$\vec{y} = \mathbf{W}\vec{x} + \vec{b}$$

1. Nonlinear activation function:

$$\vec{z} = \sigma(\mathbf{W}\vec{x} + \vec{b})$$

- Networks have many layers with many nodes

- In each layer, information from previous layer is used as input
 - This is known as the **forward pass** in a network

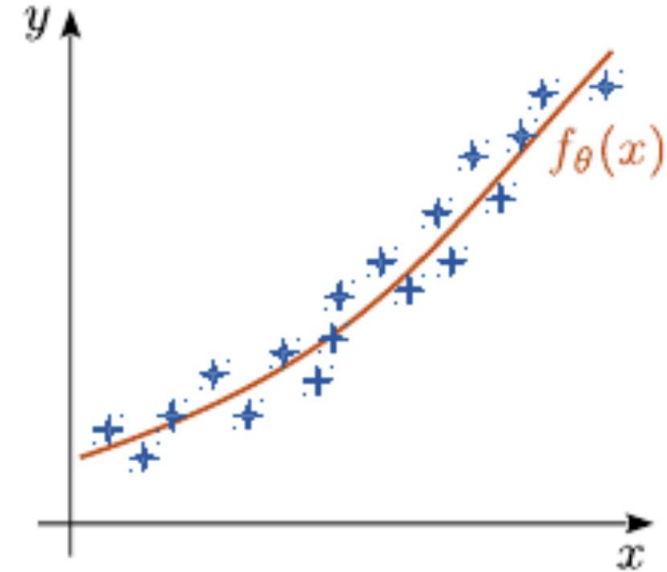


Building blocks of neural networks

- Network predictions

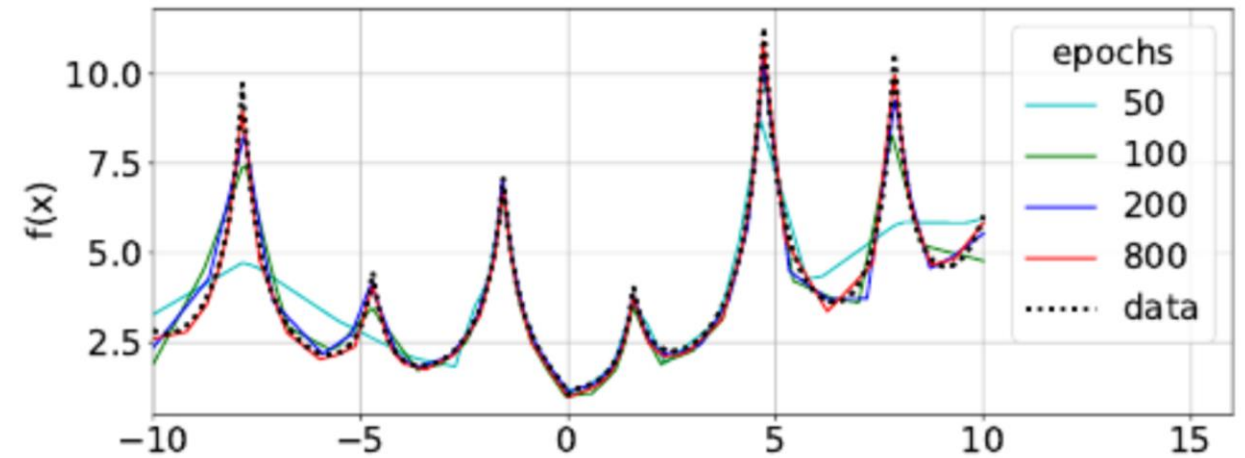
- Regression

$$z = f(x), \quad \{x, z\} \in R$$



Building blocks of neural networks

- Network predictions
 - Regression
$$z = f(x), \quad \{x, z\} \in R$$
- Example: function interpolation



Building blocks of neural networks

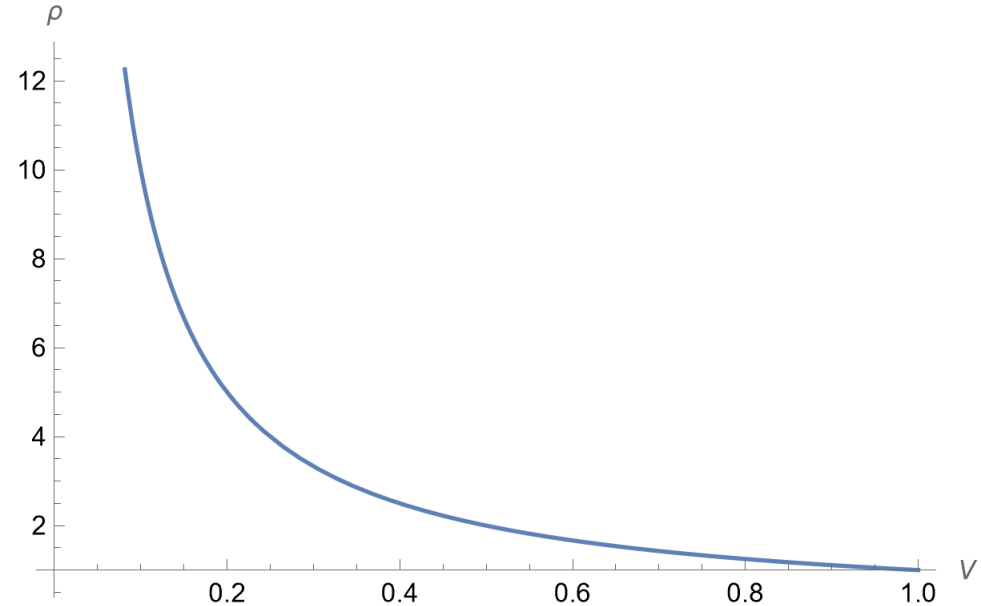
- Network predictions

- Regression

$$z = f(x), \quad \{x, z\} \in R$$

- Example: mass density

$$\rho = \frac{m}{V}$$



Building blocks of neural networks

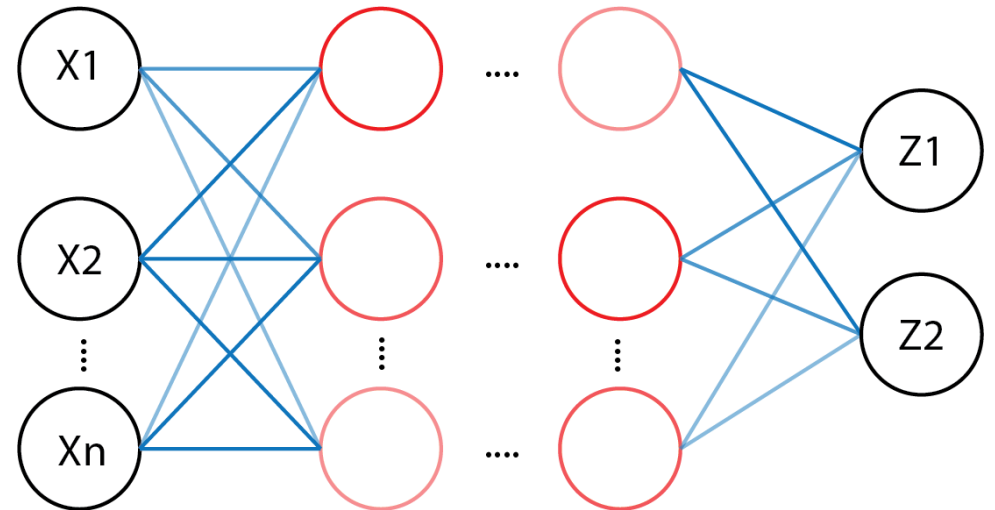
- Network predictions

- Regression

$$\vec{z} = f(\vec{x}), \quad \{x_i, z_i\} \in R$$

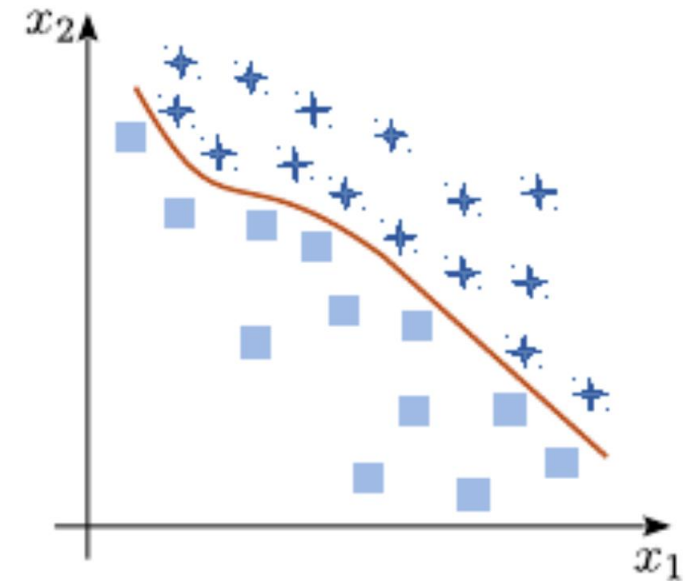
- Example: earthquake epicenter

- Inputs: time signals from each station
 - Two outputs:
 - Latitude of epicenter
 - Longitude of epicenter



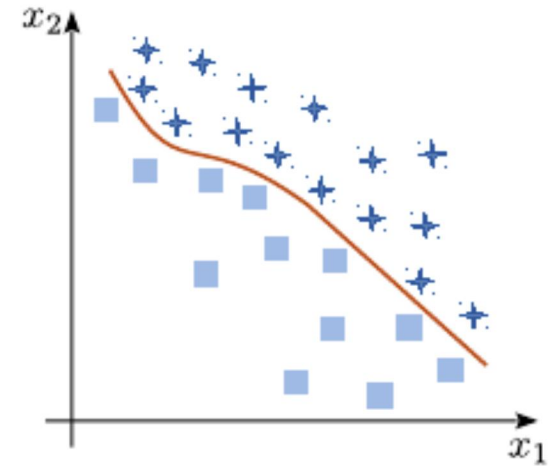
Building blocks of neural networks

- Network predictions
 - Classification:
 - Separate data into two or more categories
 - Required: separation line in the feature space



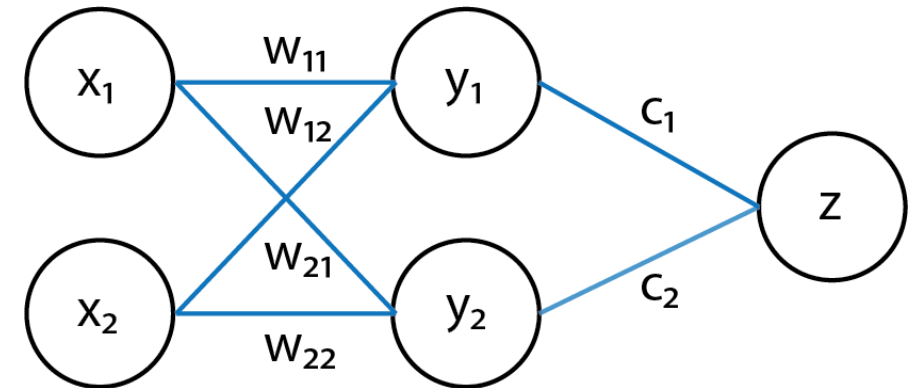
Building blocks of neural networks

- Network predictions
 - Classification:
 - Separate data into two or more categories
 - Required: separation line in the feature space



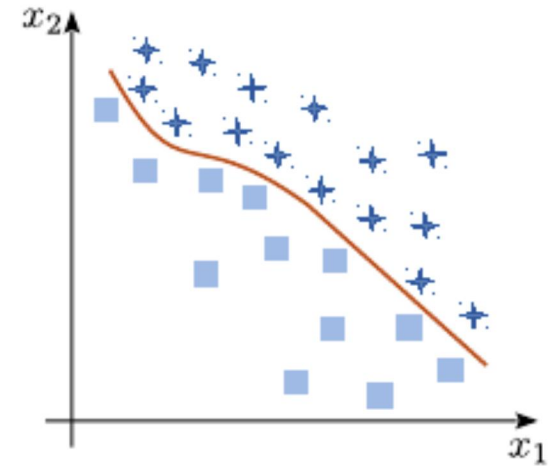
- Example:
 - separate signal from background noise
 - Trick: output is subject to a sigmoid function

$$z' = \sigma(z), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$
$$z' \in [0,1]$$



Building blocks of neural networks

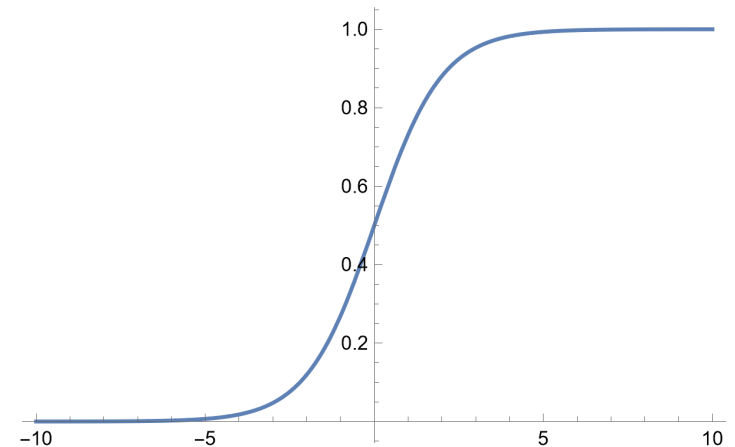
- Network predictions
 - Classification:
 - Separate data into two or more categories
 - Required: separation line in the feature space



- Example:
 - separate signal from background noise
 - Trick: output is subject to a sigmoid function

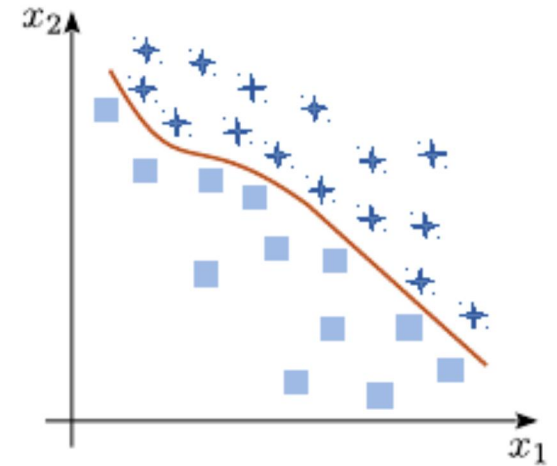
$$z' = \sigma(z), \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

$$z' \in [0, 1]$$



Building blocks of neural networks

- Network predictions
 - Classification:
 - Separate data into two or more categories
 - Required: separation line in the feature space
 - Multi-class classification
 - Network can be extended to multi-variable outputs
 - Example: classification of images into categories such as animals, vehicles, buildings



Building blocks of neural networks

- Universal approximation theorem
 - Question: can a neural network represent any function/mapping in physics?
 - *“A feed-forward network with linear output and at least one hidden layer with a finite number of nodes can approximate any of the above functions to arbitrary precision.”*
 - Theorem developed for continuous functions on closed and bounded subsets of the Euclidean space
 - Which includes all typical functions relevant to physics!

K. Hornik et al., Multilayer feedforward networks are universal approximators, Neural Networks 2 (1989)

G. Cybenko, Approximation by superpositions of a sigmoidal function, Mathematics of Control, Signals and Systems (1989)

Building blocks of neural networks

- Universal approximation theorem
 - Question: can a neural network represent any function/mapping in physics?
 - *“A feed-forward network with linear output and at least one hidden layer with a finite number of nodes can approximate any of the above functions to arbitrary precision.”*
 - The theorem has been confirmed for multilayer networks with a limited number of nodes and various activation functions
 - Learnability of the function remains an open question

K. Hornik, Approximation capabilities of multilayer feedforward networks, Neural Networks 4 (1991)

B. Hanin et al., Approximating continuous functions by ReLU Nets of minimal width, Arxiv (2017)

Z. Lu et al., The expressive power of neural networks: A view from the width, in I. Guyon et al., Advances in Neural Information Processing Systems (2017)

Summary

- Scientific questions can be answered from data using ML algorithms
- ML algorithms are trained using data by minimizing an objective function
- Physicists' tasks are to develop and improve DNN architectures suitable for solving particular tasks

Summary

- At each node of an DNN, two operations are performed:
 - Linear transformation with displacement (affine mapping)
 - Nonlinear transformation (activation function)
- A DNN can be used for
 - Regression: high-dimensional function approximation
 - Classification: classify objects into m categories
- Universal approximation theorem:
 - *“Feed-forward networks with linear output and at least one hidden layer with a finite number of nodes can approximate functions of interest in physics to arbitrary precision”*