

# Deep Learning for Physicists

Lecture #2: Training neural networks

Kosmas Kepesidis

# Some of things covered in the previous lecture...

- DNNs are trained using data by minimizing an **objective function**
- At each node of an DNN, two operations are performed:
  - Linear transformation with displacement (affine mapping)
  - Nonlinear transformation (activation function)
- Most common tasks where DNNs are used:
  - Regression: high-dimensional function approximation
  - Classification: classify objects into  $m$  categories

# Outline

- Basics of numerical optimization
- Optimization of network parameters

# Basics of numerical optimization

# Basics of numerical optimization

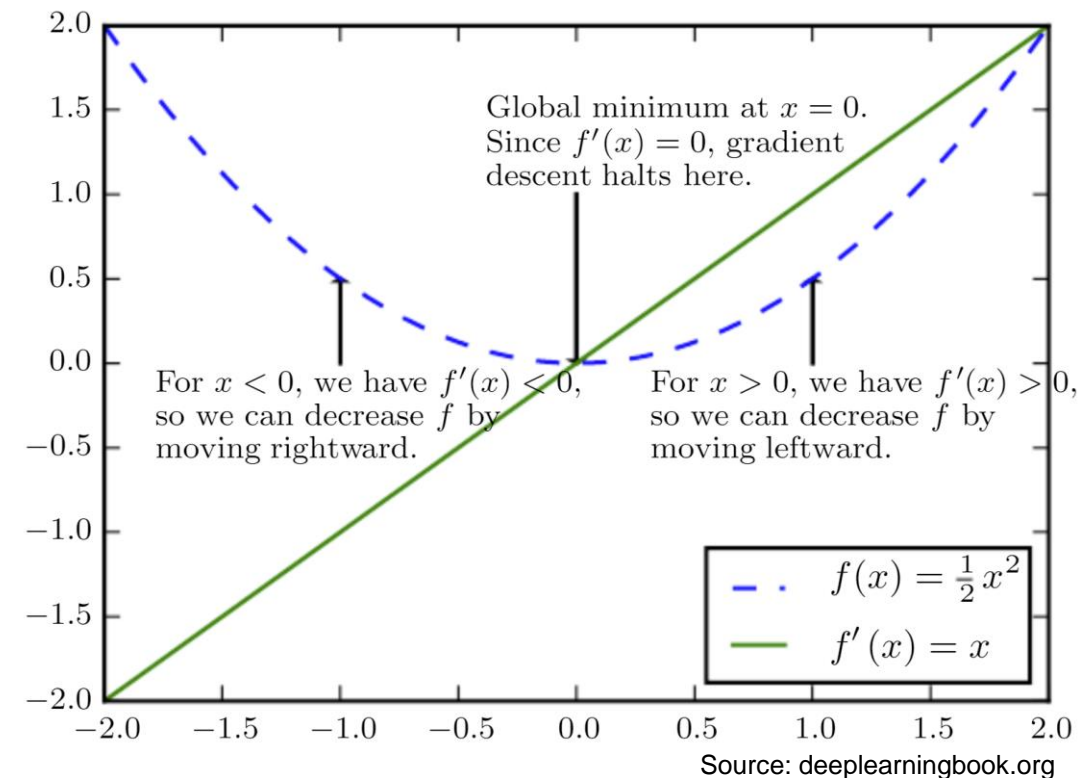
- Gradient-based optimization
  - Training deep-learning algorithms corresponds to an optimization problem
  - This is the task of minimizing (or maximizing) a certain function  $f(\mathbf{x})$  (or  $-f(\mathbf{x})$ )
    - Typically functions with multiple input and a scalar output  $f: \mathbb{R}^n \rightarrow \mathbb{R}$
  - Such a function is called **objective function**
    - Other names used: **criterion, cost function, loss function, error function**
  - Minimization using a method known as **gradient descent**

# Basics of numerical optimization

- Gradient-based optimization

- Illustration of **gradient descent** in 1D convex case

- Use first derivative  $f'(x)$  for minimization
- Look for the **global minimum**

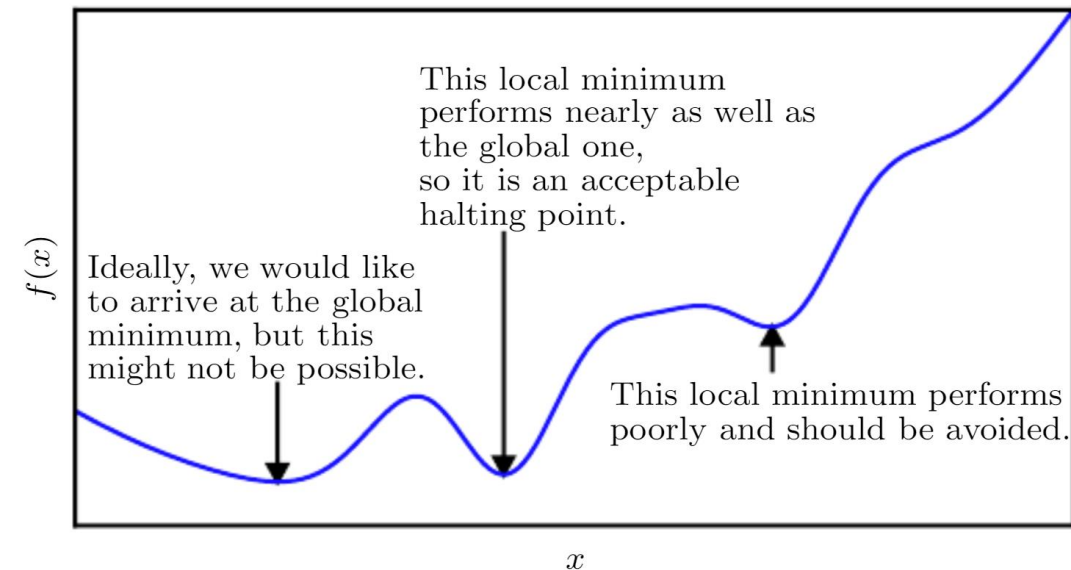


# Basics of numerical optimization

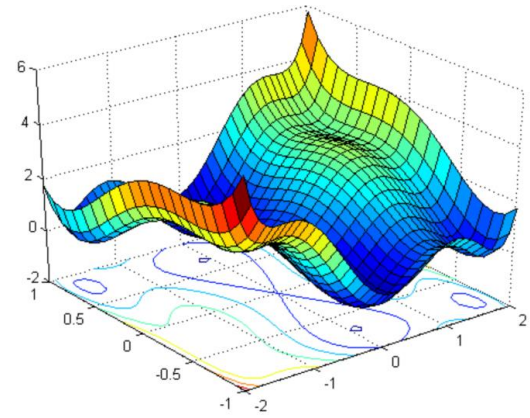
- Gradient-based optimization

- Issues in more realistic cases

- There are local **minima/maxima** and **saddle points** (critical points  $f'(x) = 0$ )
- Finding **global minimum** turns out to be a difficult job
- Satisfying solutions of a “good” local minimum
- Avoid poor-performing “bad” local minima



# Basics of numerical optimization



Source: <https://algorithmia.com>

- Gradient-based optimization
  - Typically, functions with multiple input and a scalar output are used  $f: \mathbb{R}^n \rightarrow \mathbb{R}$
  - In multiple dimensions, the location of minima (critical points in general) is defined the partial derivative in each direction needs to be zero  $\nabla_{\mathbf{x}} f(\mathbf{x}) = 0$
  - To minimize  $f$  one needs to find the direction in which  $f$  decreases the fastest – this is done using the **directional derivative** (slope of  $f$  in direction of unit vector  $\mathbf{u}$ )
    - See <https://www.deeplearningbook.org/contents/numerical.html> for more details
  - $f$  can be decreased by moving in the opposite direction of the negative gradient – a method known as the **steepest descend** or **gradient descent**:  $\mathbf{x}' = \mathbf{x} - \alpha \nabla_{\mathbf{x}} f(\mathbf{x})$ 
    - $\alpha$  is known as the **learning rate**



# Basics of numerical optimization

- Gradient-based optimization
  - Performance of **gradient descent** depends on the differences in scale between input features of the data
  - Large differences in scales can lead to very long convergence times

Optimization of network parameters

# Optimization of network parameters

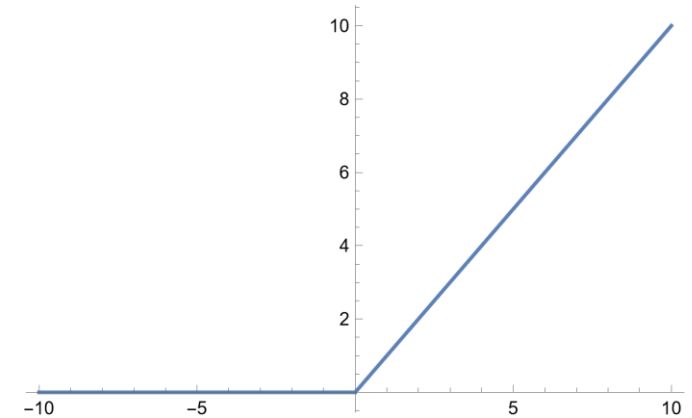
- Input data preprocessing and numerical stability
  - Numerical stability is crucial for successful training of neural networks
  - If the ranges of the input data vary by many orders of magnitude and are different for the individual components, adjusting network parameters becomes difficult
  - Careful preparation and pre-processing of the input data are essential for successful model building

# Optimization of network parameters

- Input data preprocessing and numerical stability

## ➤ Zero-centering:

- Gradient of widely-used ReLU activation changes dramatically close to zero
- This can affect the importance of results of the affine mappings  $\vec{y} = \mathbf{W} \vec{x} + \vec{b}$
- Solution: subtract mean values  $x_i \rightarrow x_i - \langle x_i \rangle$



# Optimization of network parameters

- Input data preprocessing and numerical stability

- Scaling of data

- If the values of different observables  $x_i$  are of different orders of magnitude, larger values might be considered more “significant”, when training a neural network

- Solution: standard scaling  $x_i \rightarrow \frac{x_i - \langle x_i \rangle}{\sigma_i}$

- If values in fixed interval  $[x_{i,min}, x_{i,max}]$ :  $x_i \rightarrow 2 \frac{x_i - x_{i,min}}{x_{i,max} - x_{i,min}} - 1$

# Optimization of network parameters

- Input data preprocessing and numerical stability

- Logarithmic scale

- For variables whose values fluctuate, it might be useful to redistribute values by transforming them using the logarithm (or exponential)

$$x_i \rightarrow \log(x_i)$$

$$x_i \rightarrow e^{-x_i}$$

# Optimization of network parameters

- Input data preprocessing and numerical stability

➤ Decorrelation – reduce redundancy

- Identify colinear variables and keep only one:  $x_i = c * x_j$
- Identify variables with a known functional form between them:  $x_i = f(x_j)$

# Optimization of network parameters

- Input data preprocessing and numerical stability

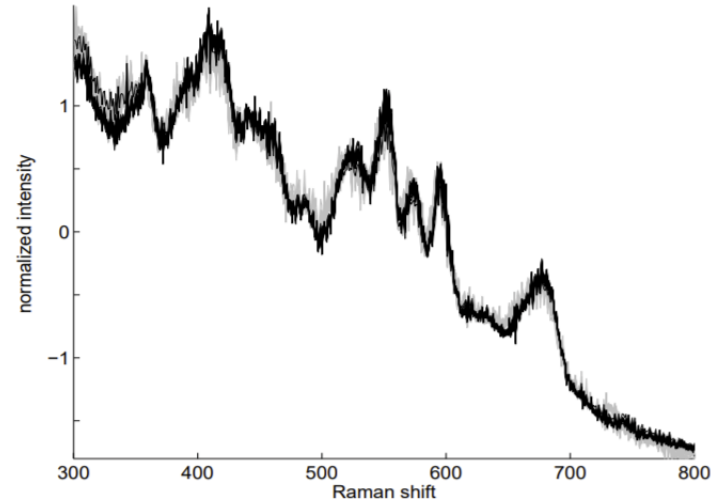
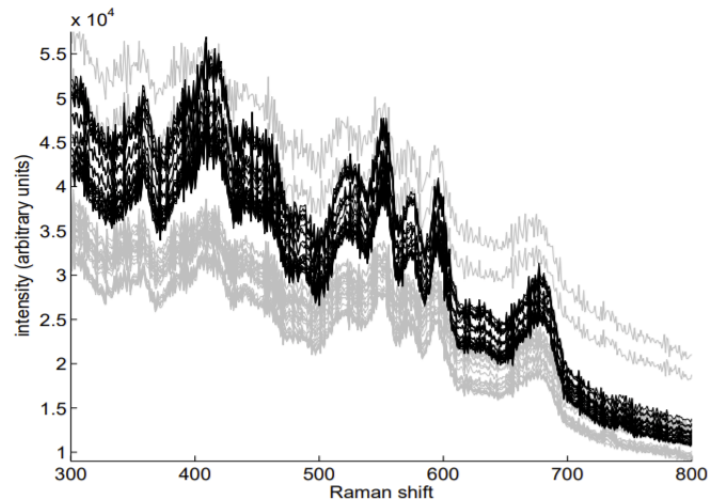
- Normalization (global)

- Series of measurements containing same quantity can be normalized with respect to a global value such as global average, maximum etc.
    - Examples:
      - Pixel intensity of an image
      - Absorbance in infrared (IR) spectroscopy



# Optimization of network parameters

- Input data preprocessing and numerical stability
  - Real example: Raman spectroscopy



# Optimization of network parameters

- Input data preprocessing and numerical stability
  - Thought experiment with seismic stations
    - 50 measuring stations record seismic movements as function of time
    - 100k earthquake events collected, in time span of 5 years
    - Data includes the times  $t_{p,i}$  and  $t_{s,i}$  of maximum amplitudes  $A_{s,i}$  and  $A_{p,i}$  of compression and shear waves respectively
  - Scaling the time information:
    - $t_{p,i}$  and  $t_{s,i}$  are measured in seconds after Unix timestamp (01.01.1970 at 00:00)
    - Such big numbers affect network performance
    - Solution: scaling using first and last recordings

$$t_{p,i} \rightarrow 2 \frac{t_{p,i} - t_{p,i,min}}{t_{p,i,max} - t_{p,i,min}} - 1$$

# Optimization of network parameters

- Input data preprocessing and numerical stability
  - Thought experiment with seismic stations
    - 50 measuring stations record seismic movements as function of time
    - 100k earthquake events collected, in time span of 5 years
    - Data includes the times  $t_{p,i}$  and  $t_{s,i}$  of maximum amplitudes  $A_{s,i}$  and  $A_{p,i}$  of compression and shear waves respectively
  - Log-transform amplitude:
    - Amplitude values are assumed to be in an interval  $[0,1000]$  – with the majority close to zero
    - Log-transform is appropriate:  $A_{s,i} \rightarrow \log(A_{s,i})$

# Optimization of network parameters

- After data preprocessing:
  - Training of DNNs is an iterative process
  - Available data is used repeatedly and as efficiently as possible
  - Two terms are associated with this procedure: **epoch** and **minibatch**

# Optimization of network parameters

- **Epoch**

- One epoch of a network training denotes the one-time, complete use of all training data

# Optimization of network parameters

- **Minibatch**

- The parameters of a network are iteratively optimized in many small steps

- Using all training data in each step would be time-consuming and highly inefficient
    - Use instead a randomly selected sample of the training data in each iterative step
    - Size of the batch, in powers of 2 ( $k = 2^m$ ) to be chosen depends on the problem to be solved
    - Tradeoff: computing **costs**  $\sim k$  to be balanced by **precision**  $\sim \frac{1}{\sqrt{k}}$

# Optimization of network parameters

- Network initialization

➤ While  $b_i$  are set to zero, initial weights  $\mathbf{W}$  in mappings  $\vec{y} = \mathbf{W}\vec{x} + \vec{b}$  are usually selected to be random numbers taken from

1. uniform  $[-s, s]$
2. or standard normal  $N(\mu = 0, \sigma)$  distribution

➤ For normally-distributed random weights, standard deviation should be optimally scaled by the numbers of input/output neurons ( $n_{in}$ ,  $n_{out}$ )

- For tanh activation, recommend:  $\sigma^2 = \frac{2}{n_{in} + n_{out}}$  (Glorot normal or Xavier normal initialization)
- For ReLU activation, recommend:  $\sigma^2 = \frac{2}{n_{in}}$

# Optimization of network parameters

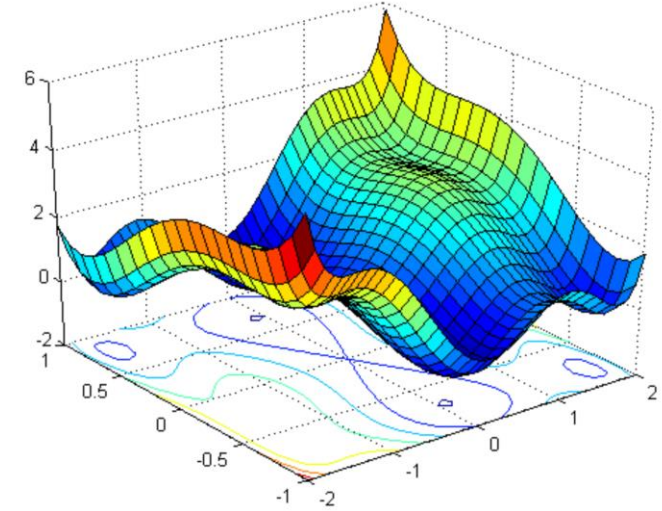
- Network initialization

- While  $b_i$  are set to zero, initial weights  $\mathbf{W}$  in mappings  $\vec{y} = \mathbf{W}\vec{x} + \vec{b}$  are usually selected to be random numbers taken from
  1. uniform  $[-s, s]$
  2. or standard normal  $N(\mu = 0, \sigma)$  distribution
- In the case of uniformly-distributed random weights, to recover the same spread of weights  $\mathbf{W}$  as for the Gaussian initialization, use  $s = 3 \times \sigma$



# Optimization of network parameters

- **Objective function** (also known as **cost**, **loss** function) is a measure for evaluating networks predictions
  - It depends on all (to-be-optimized) parameters of the neural network
  - Approach a “good” local minimum on the hyper-plane of all parameters
  - The “goodness” of the local minimum is assessed by evaluation of the network predictions



Source: <https://algorithmia.com>

# Optimization of network parameters

- Objective function for **regression**:

- Distance measure between predictions  $f(x_i)$  and target values  $y(x_i)$ , where  $i$  runs over data points

- **Mean absolute error (MAE)** – *Manhattan norm*  $\mathcal{L} = \frac{1}{k} \sum_{i=1}^k |f(x_i) - y(x_i)|$

- **Mean squared error (MSE)**  $\mathcal{L} = \frac{1}{k} \sum_{i=1}^k [f(x_i) - y(x_i)]^2$

- **Root mean squared error (RMSE)** – *Euclidean norm*  $\mathcal{L} = \sqrt{\frac{1}{k} \sum_{i=1}^k [f(x_i) - y(x_i)]^2}$

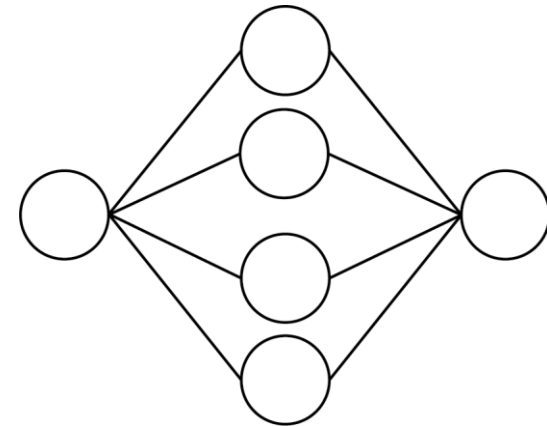
# Optimization of network parameters

- Objective function for **regression**:

- Distance measure between predictions  $f(x)$  and target values  $y(x)$

- **Mean squared error (MSE)**  $\mathcal{L} = \frac{1}{k} \sum_{i=1}^k [f(x_i) - y(x_i)]^2$

- For 1D input
    - Does not account for multiple observables



# Optimization of network parameters

- Objective function for **regression**:

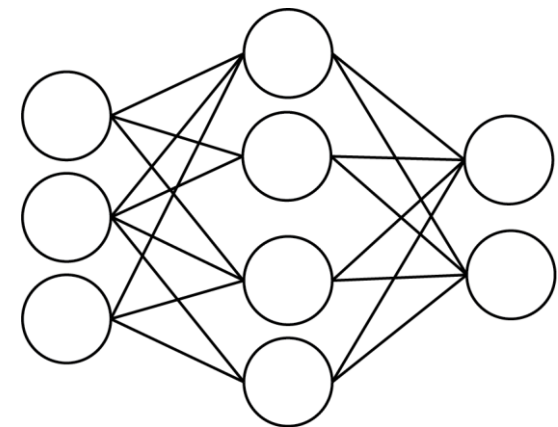
- distance measure between predictions  $f(x)$  and target values  $y(x)$

- **Mean squared error (MSE)**

$$\mathcal{L} = \frac{1}{k} \sum_{i=1}^k \sum_{j=1}^m [f(\vec{x}_i) - y(\vec{x}_i)]^2$$

- Extended MSE to  **$n$ -dimensional input** and  **$m$ -dimensional output**

- $\vec{x}_i = \begin{pmatrix} x_{i,1} \\ \vdots \\ x_{i,n-1} \\ x_{i,n} \end{pmatrix}$



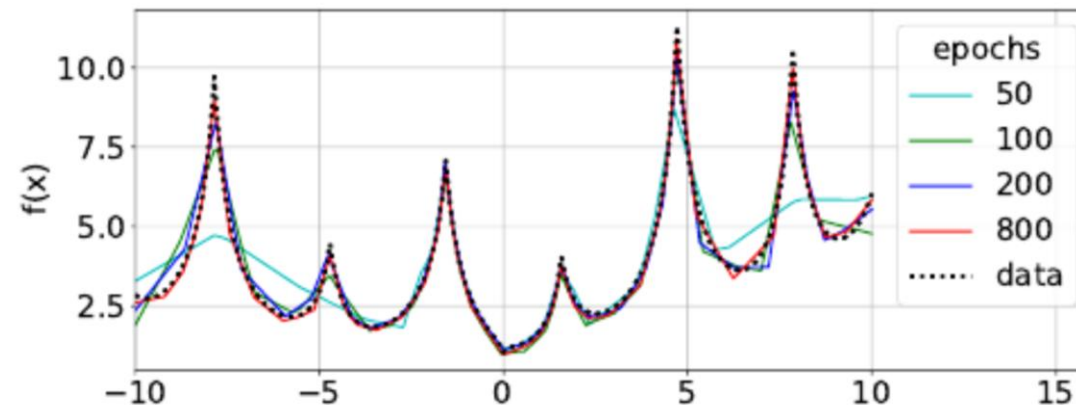
# Optimization of network parameters

- Objective function for **regression**:
  - distance measure between predictions  $f(x)$  and target values  $y(x)$

- **Mean squared error (MSE)**

$$\mathcal{L} = \frac{1}{k} \sum_{i=1}^k \sum_{j=1}^m [f(\vec{x}_i) - y(\vec{x}_i)]^2$$

- Example: Function interpolation



Source: M. Erdmann et al., Deep Learning for Physics Research

# Optimization of network parameters

- Objective function for **classification**
  - Objective function needs to process probabilities
  - Solution comes from statistical mechanics:
    - **Boltzmann entropy:**  $S = k_B \log W$ 
      - $k_B$ : Boltzmann constant
      - $W$ : number of all equally-probable configurations in a system
    - Probability:  $p = 1/W$ , therefore  $S = -k_B \log(p)$
    - Generalize to non-uniform probability distributions:
      - $S = -k_B \langle \log(p_j) \rangle = -k_B \sum_j p_j \log(p_j)$

# Optimization of network parameters

- Objective function for **classification**

➤ **Shannon's entropy**: measure of ignorance in information theory

- $S = -k_S \sum_j p_j \log(p_j)$ , where  $k_S = \frac{1}{\log(2)}$
- $S = -\sum_j p_j \log_2(p_j)$
- It is a measure of entropy in bits

# Optimization of network parameters

- Objective function for **classification**

➤ **Cross-entropy**: distinguishes between true probabilities ( $p_j$ ) and estimated probabilities ( $q_j$ )

$$H = - \sum_j p_j \log(q_j)$$

- Interesting and useful property:  $H$  becomes minimal for  $p_j = q_j$

➤ Objective function based on cross entropy:

$$\mathcal{L} = - \frac{1}{k} \sum_{i=1}^k \left[ \sum_{j=1}^m p_j \log(q_j) \right]_i$$



# Optimization of network parameters

- Objective function for **classification**

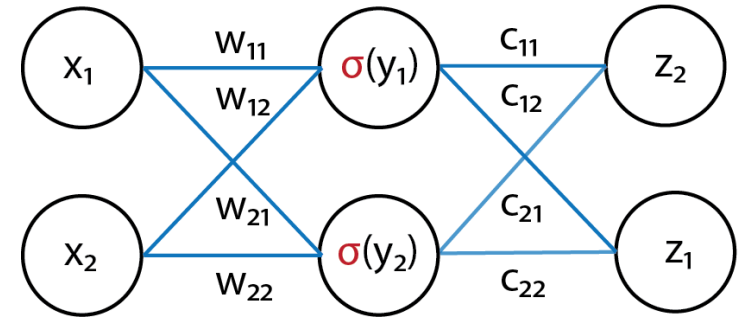
➤ For binary classification problem i.e.  $m = 2$  and  $p_1 + p_2 = 1$ , cross-entropy is given by:

$$\mathcal{L} = -\frac{1}{k} \sum_{i=1}^k [p_1 \log(q_1) + (1 - p_1) \log(1 - q_1)]_i$$

# Optimization of network parameters

- Objective function for **classification**

➤ Example: separating signal from noise



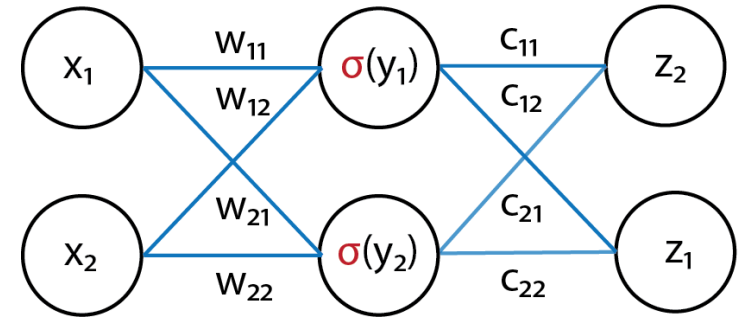
- NN with two output nodes: 
$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} c_{11} & c_{12} \\ c_{21} & c_{22} \end{pmatrix} \begin{pmatrix} \sigma(y_1) \\ \sigma(y_2) \end{pmatrix} + \begin{pmatrix} d_1 \\ d_2 \end{pmatrix}$$
- Activation function: hyperbolic tangent  $\sigma(y) = \tanh(y)$
- Want to transform results of output layer into probabilities, i.e.  $\hat{z}_1 + \hat{z}_2 = 1$
- This is done using the **softmax** function: 
$$\vec{q} \equiv \begin{pmatrix} \hat{z}_1 \\ \hat{z}_2 \end{pmatrix} = \frac{1}{e^{z_1} + e^{z_2}} \begin{pmatrix} e^{z_1} \\ e^{z_2} \end{pmatrix}$$

# Optimization of network parameters

- Objective function for **classification**

➤ Example: separating signal from noise

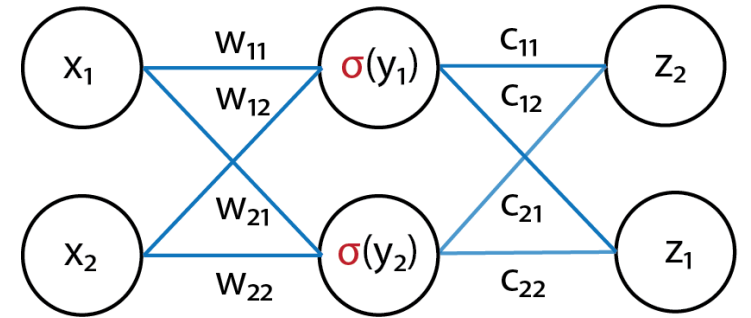
- NN with two output nodes: 
$$\begin{pmatrix} z_1 \\ z_2 \end{pmatrix} = \begin{pmatrix} c_{11} & c_{12} \\ c_{21} & c_{22} \end{pmatrix} \begin{pmatrix} \sigma(y_1) \\ \sigma(y_2) \end{pmatrix} + \begin{pmatrix} d_1 \\ d_2 \end{pmatrix}$$
- Activation function: hyperbolic tangent  $\sigma(y) = \tanh(y)$
- Transform results of output layer into probabilities, i.e.  $\hat{z}_1 + \hat{z}_2 = 1$
- This is done using the **softmax** function: 
$$\vec{q} \equiv \begin{pmatrix} \hat{z}_1 \\ \hat{z}_2 \end{pmatrix} = \frac{1}{e^{z_1} + e^{z_2}} \begin{pmatrix} e^{z_1} \\ e^{z_2} \end{pmatrix}$$
- Signal:  $\hat{z}_1 > \hat{z}_2$ , noise:  $\hat{z}_1 < \hat{z}_2$



# Optimization of network parameters

- Objective function for **classification**

➤ Example: separating signal from noise



➤ True values encoded in **one-hot-encoding**:  $\vec{p}^T = (p_1 \ p_2) = \begin{cases} (1 & 0): \text{Signal} \\ (0 & 1): \text{Noise} \end{cases}$

➤ For training the network we use the cross-entropy for a mini-batch of size  $k$  with data and  $m = 2$  classes

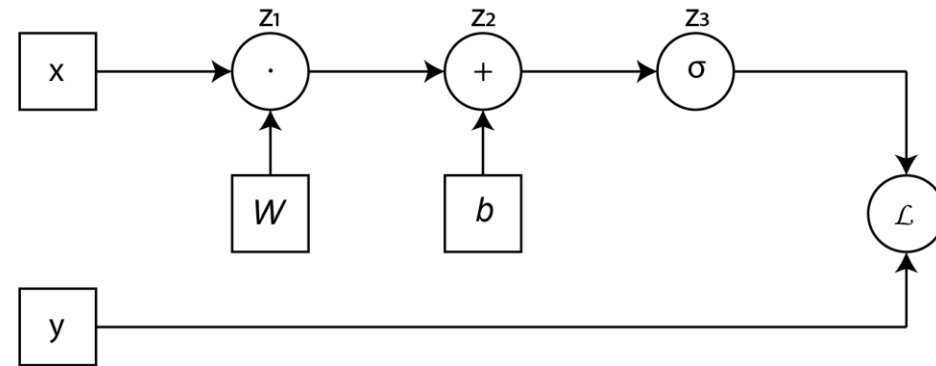
$$\mathcal{L} = -\frac{1}{k} \sum_{i=1}^k \left[ \sum_{j=1}^{m=2} p_j(\vec{x}_i) \log(q_j(\vec{x}_i)) \right] = -\frac{1}{k} \sum_{i=1}^k \vec{p}^T(\vec{x}_i) \log(\vec{q}^T(\vec{x}_i))$$

# Optimization of network parameters

- Numerical optimization is achieved using **gradient descent**
  - it is based on gradients of objective function with respect to the network parameters  $\vec{W}$  and  $\vec{b}$
  - The gradients are obtained using a method called **backpropagation**

➤ Example with a single node:

$$\begin{aligned} z_1 &= Wx & z_3 &= \sigma(z_2) \\ z_2 &= z_1 + b & \mathcal{L} &= (y - z_3)^2 \end{aligned}$$



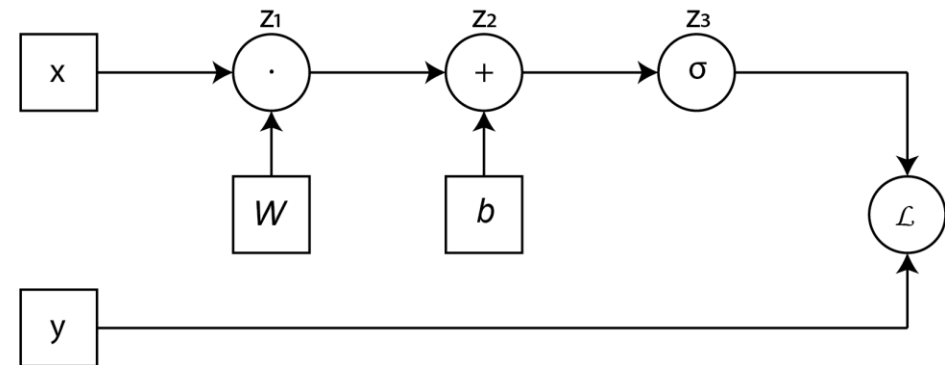
Question: How can parameters be changed to achieve minimal values of  $\mathcal{L}$  ?

# Optimization of network parameters

- Using the chain rule, calculate partial derivatives of  $\mathcal{L}$  with respect to the  $W$  and  $b$ :

$$\frac{\partial \mathcal{L}}{\partial W} = \frac{\partial \mathcal{L}}{\partial z_3} \frac{\partial z_3}{\partial z_2} \frac{\partial z_2}{\partial z_1} \frac{\partial z_1}{\partial W}$$

$$\frac{\partial \mathcal{L}}{\partial b} = \frac{\partial \mathcal{L}}{\partial z_3} \frac{\partial z_3}{\partial z_2} \frac{\partial z_2}{\partial b}$$



- Calculation of derivative done opposite to the direction of forward pass:  
**backpropagation**

# Optimization of network parameters

- Optimizing with **stochastic gradient descent (SGD)**

➤ Gradients of objective function calculated as average over  $k$  data points of the minibatch

$$\mathbb{E} \left[ \frac{\partial \mathcal{L}}{\partial W} \right] = \frac{1}{k} \sum_{i=1}^k \frac{\partial \mathcal{L}}{\partial W} \quad \mathbb{E} \left[ \frac{\partial \mathcal{L}}{\partial b} \right] = \frac{1}{k} \sum_{i=1}^k \frac{\partial \mathcal{L}}{\partial b}$$

➤ SGD leads to greater variance

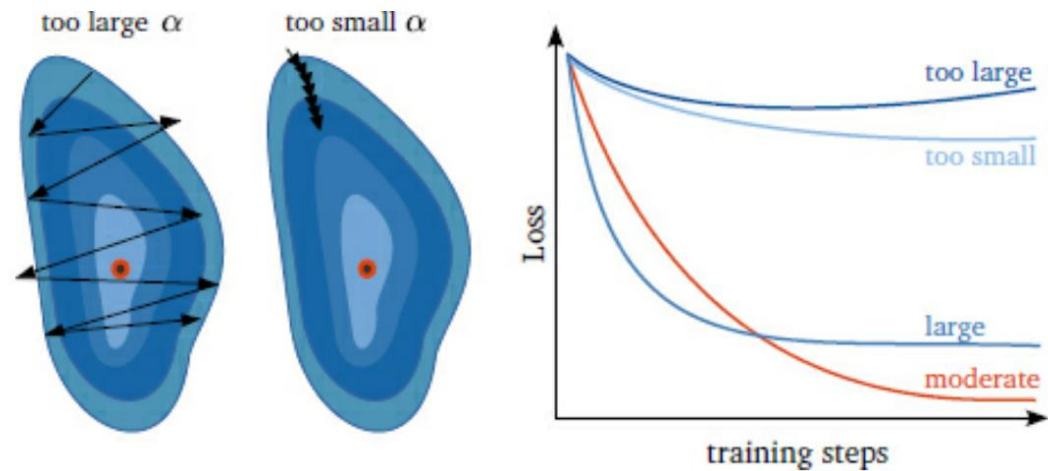
- Performing optimization in terms of minibatches enables multiple parameter updates in a single epoch
- This approach is more robust against “unwanted” local minima

# Optimization of network parameters

- SGD evaluates whether parameter  $W$  (and/or  $b$ ) needs to be increased or decreased in the next iteration step
- How large the change in the parameters will be (step size) is determined by the **learning rate:  $\alpha$**

$$W_{t+1} = W_t - \alpha \mathbb{E} \left[ \frac{\partial \mathcal{L}}{\partial W} \right]_t$$

$$b_{t+1} = b_t - \alpha \mathbb{E} \left[ \frac{\partial \mathcal{L}}{\partial b} \right]_t$$



Source: M. Erdmann et al., Deep Learning for Physics Research



# Optimization of network parameters

- SGD evaluates whether parameter  $W$  (and or  $b$ ) needs to be increased or decreased in the next iteration step
- How large the change in the parameters will be (step size) is determined by the **learning rate**:  $\alpha$
- Learning rates are varied in the training process
  - Initially, parameters are examined with larger step sizes
  - Usually in range:  $\alpha = 10^{-5} - 10^{-2}$
  - Subsequently,  $\alpha$  is gradually reduced: examining parameters for smaller and smaller step size

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

## 1. **Adagrad** (adaptive gradient): adaptive learning rates

- During optimization,  $\alpha$  reduces continuously
- Changes in  $\alpha$ , are adapted individually for each parameter
- Takes into account sum of squares of all previous gradients

$$v_t = \sum_{\tau=1}^t \left( \frac{\partial \mathcal{L}}{\partial W} \right)^2 \quad \alpha_t = \frac{\alpha}{\sqrt{v_t} + \epsilon}$$

- $\epsilon \approx 10^{-8}$  guaranties not division by zero

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

## 2. **RMSprob**: adaptive learning rates as well

- Includes decay parameter to suppress influence from gradients of older steps
- Decay parameter with typical value:  $\beta = 0.9$

$$v_t = \beta v_{t-1} + (1 - \beta) \left( \frac{\partial \mathcal{L}}{\partial W} \right)^2 \qquad \alpha_t = \frac{\alpha}{\sqrt{v_t} + \epsilon}$$

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

3. **Momentum**: considers parameters and gradients geometrically

$$\vec{\theta} = \begin{pmatrix} W \\ b \end{pmatrix} \quad \nabla \vec{\theta} = \begin{pmatrix} \frac{\partial \mathcal{L}}{\partial W} \\ \frac{\partial \mathcal{L}}{\partial b} \end{pmatrix}$$

➤ Goal the most efficient step size in the right direction in parameter space

$$\vec{u}_t = -\alpha \nabla \vec{\theta}_t \quad \vec{\theta}_{t+1} = \vec{\theta}_t + \vec{u}_t$$

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

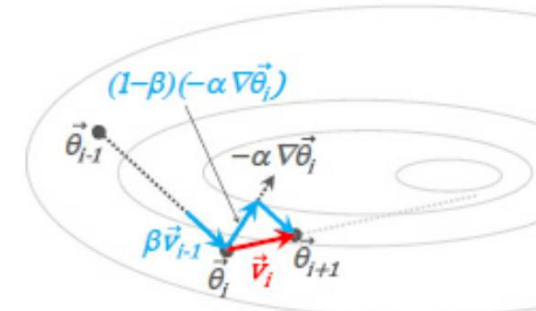
## 3. **Momentum**: considers parameters and gradients geometrically

- Essential aspect of momentum method: Stabilize direction of optimization using the history of velocity

$$\vec{u}_t = \beta \vec{u}_{t-1} + (1 - \beta) (-\alpha \nabla \vec{\theta}_t)$$

$$\vec{\theta}_{t+1} = \vec{\theta}_t + \vec{u}_t$$

- Coefficient  $\beta$  balances influence of previous velocity and its modification
  - Typical values:  $\beta = 0.5, \dots, 0.9$



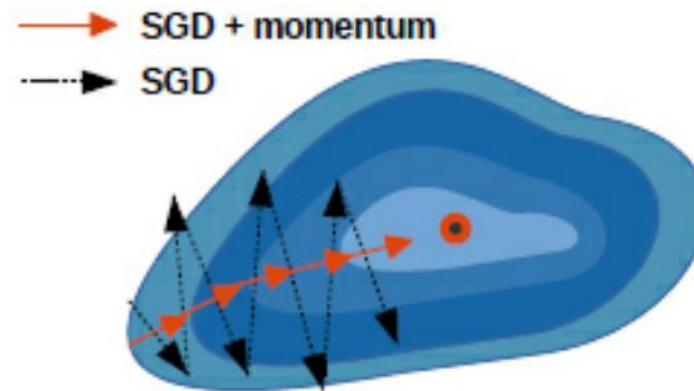
Source: M. Erdmann et al., Deep Learning for Physics Research

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

3. **Momentum**: considers parameters and gradients geometrically

➤ Leads to oscillating behavior in the parameter space – with damping!



Source: M. Erdmann et al., Deep Learning for Physics Research

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

4. **Adam** (adaptive moments): combines ideas of RMSprop and the momentum method

➤ Both gradients and their squares are subject to decay:

$$m_t = \frac{1}{1 - \gamma^t} \left[ \gamma m_{t-1} + (1 - \gamma) \frac{\partial \mathcal{L}}{\partial W} \right]$$

$$v_t = \frac{1}{1 - \beta^t} \left[ \beta v_{t-1} + (1 - \beta) \left( \frac{\partial \mathcal{L}}{\partial W} \right)^2 \right]$$

- Proposed coefficient values:  $\gamma = 0.9$ ,  $\beta = 0.999$
- Norm factors loose influence for  $t \gg 1$

# Optimization of network parameters

- Learning strategies: methods for accelerating network training!

4. **Adam** (adaptive moments): combines ideas of RMSprob and the momentum method

- $m_t$ : scales direction of next step in parameter space
- $v_t$ : adapts learning rates

$$\vec{u}_t = -\alpha \frac{m_t}{\sqrt{v_t} + \epsilon} \quad \vec{\theta}_{t+1} = \vec{\theta}_t + \vec{u}_t$$

- Adam is efficient & robust: good first choice as optimizer!



# Summary

- Data must be appropriately preprocessed before inserted to neural network: especially ***scaling*** reduces high fluctuations in data
- During training, the data set is used multiple times – each time is called ***Epoch***
- Parameter optimization is done in smaller steps using only samples of the data set called ***minibatches***
- Weight coefficients are initialized using random numbers following a distribution, shoes variance depends on the number of nodes in a layer and the ***activation function***
- Common ***objective functions*** for regression are ***MAE, MSE, RMAE*** and for classification we use ***cross-entropy***
- With the help of the chain rule of partial derivatives (***backpropagation***), ***stochastic gradient descent*** minimizes ***objective function***
- ***Learning rate*** corresponds to the steps size of the optimization procedure
- Different optimization strategies can be chosen – during optimization the learning rate is dynamically adapt for efficiency and robust results