

# Deep Learning for Physicists

Lecture #3: Practical methodology

Kosmas Kepesidis

# What was covered in the previous lecture...

- During training, the data set is used multiple times – each time is called ***Epoch***
- Parameter optimization is done in smaller steps using only samples of the data set called ***minibatches***
- Weight coefficients are initialized using random numbers following a distribution (normal or uniform)
- Common ***objective functions*** for regression are ***MAE, MSE, RMAE*** and for classification we use ***cross-entropy***
- With the help of the chain rule of partial derivatives (***backpropagation***), ***stochastic gradient descent*** minimizes ***objective function***
- ***Learning rate*** corresponds to the steps size of the optimization procedure

# Outline

- Practical methodology
  - Criteria for model training
  - Train, validation and test data sets
  - Monitoring
  - Regularization
  - Hyperparameters

# Criteria for model training

# Criteria for model training

- Objective function for regression:

➤ Distance measure between predictions  $f(x_i)$  and target values  $y(x_i)$ , where  $i$  runs over data points

- Mean absolute error (MAE) – *Manhattan norm*  $\mathcal{L} = \frac{1}{k} \sum_{i=1}^k |f(x_i) - y(x_i)|$

- Mean squared error (MSE)  $\mathcal{L} = \frac{1}{k} \sum_{i=1}^k [f(x_i) - y(x_i)]^2$

- Root mean squared error (RMSE) – *Euclidean norm*  $\mathcal{L} = \sqrt{\frac{1}{k} \sum_{i=1}^k [f(x_i) - y(x_i)]^2}$


# Criteria for model training

- Model generalization: regression case
  - $n$  data points  $(x_i, y_i)$  generated by a stochastic process:  $y_i = g(x_i) + \varepsilon$ 
    - $g(x_i)$ : probability distribution
    - $\varepsilon$ : noise term following standard normal distribution  $\sim N(0, \sigma)$
  - A network output  $f(x)$  is optimized using  $MSE = \langle (f - y)^2 \rangle$
  - Bias-variance relation:  $MSE = B[\langle f \rangle, \langle g \rangle]^2 + V[f] + \sigma^2$ 
    - $B[\langle f \rangle, \langle g \rangle]$ : bias term, i.e. displacement of the expectation value of the network prediction versus true value
    - $V[f]$ : variance of network predictions
    - $\sigma^2$ : noise variance – irreducible uncertainty

# Criteria for model training

- Model generalization: regression case

➤  $MSE = B[\langle f \rangle, \langle g \rangle]^2 + V[f] + \sigma^2$




- Models with low bias and low variance have **generalization** capability

# Criteria for model training

- Model generalization

➤ Bias-variance tradeoff:






# Criteria for model training

- **Overfitting/underfitting**

- Network model should capture complexity of true distribution without being more complicated than it should
- Oversimplified models can cause underfitting
- Too complicated models can potentially capture fluctuations in the data (overfitting)

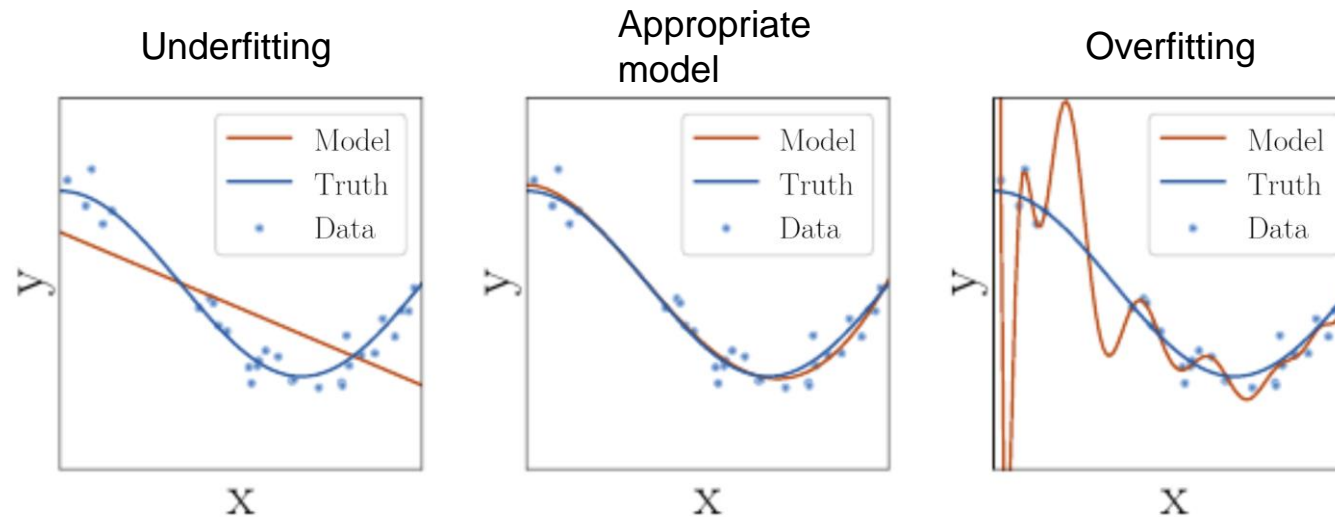


- DNNs have many parameters and are prone to overfitting
- training should be monitored!

# Criteria for model training


- **Overfitting/underfitting**

- Network model should capture complexity of true distribution without being more complicated than it should



# Criteria for model training

- **Overfitting/underfitting**



# Criteria for model training

- **Overfitting/underfitting**

- Network model should capture complexity of true distribution without being more complicated than it should

- Learner finds a pattern in the data that is not actually true in the real world: overfitting
      - Happens when you have too many hypotheses and not enough data to tell them apart
      - The more data, the more “bad” hypotheses are eliminated
    - If the hypothesis space is not constrained, there may never be enough data
      - There is often a parameter that allows you to constrain (regularize) the learner

# Criteria for model training

- **Overfitting/underfitting**

- Why does overfitting occur and how we avoid it:

- Happens when you have too many hypotheses and not enough data to tell them apart  
→ collect more data
    - Data is noisy  
→ collect better data (reduce noise)
    - Models are too complex → use less complex models
    - Aggressive loss optimization → optimize less

# Criteria for model training

- Evaluation metrics for classification tasks

➤ **Confusion matrix** for binary classification

		Predicted class	
		Positive	Negative
Actual class	Positive	True positives (TP)	False negatives (FN)
	Negative	False positives (FP)	True negatives (TN)

# Criteria for model training

- Evaluation metrics for classification tasks

- **Confusion matrix** for binary classification

- Example: spam emails

	spam (predicted)	not_spam (predicted)
spam (actual)	23 (TP)	1 (FN)
not_spam (actual)	12 (FP)	556 (TN)

- Question: how accurate is this classifier?

# Criteria for model training

- Evaluation metrics for classification tasks

- **Confusion matrix** for binary classification

- **Accuracy:**

- how many predictions are correct?

$$accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

- Overall measure of classification performance
      - Can be problematic if classes are imbalanced

		Predicted class	
		Positive	Negative
Actual class	Positive	True positives (TP)	False negatives (FN)
	Negative	False positives (FP)	True negatives (TN)



# Criteria for model training

- Evaluation metrics for classification tasks

- **Confusion matrix** for binary classification

- **Precision:**

- how many of the positive predictions are correct?

$$precision = \frac{TP}{TP + FP}$$

- **Recall (Sensitivity):**

- how many of positives were correctly identified?

$$recall = \frac{TP}{TP + FN}$$

		Predicted class	
		Positive	Negative
Actual class	Positive	True positives (TP)	False negatives (FN)
	Negative	False positives (FP)	True negatives (TN)

# Criteria for model training

- Evaluation metrics for classification tasks

- **Confusion matrix** for binary classification

- **Precision - Recall:**

- Can be combined into a single measure: *F1*-score

$$F1 = 2 \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{TP}{TP + \frac{FN + FP}{2}}$$

- *F1*-score favors classifiers with balanced precision and recall

		Predicted class	
		Positive	Negative
Actual class	Positive	True positives (TP)	False negatives (FN)
	Negative	False positives (FP)	True negatives (TN)

# Criteria for model training

- Evaluation metrics for classification tasks

- **Precision/recall tradeoff:**


One typically has to choose between high precision or high recall – increasing one reduces the other

- Consider classification task: class  $A$  or class  $\bar{A}$  (i.e. *not* class  $A$ )

- A fixed threshold for the classification can be used:  $p_A > 0.5$

- Alternatively,  $p_A > v$ ,  $v \in [0,1]$

- Threshold  $v$  is referred to as **operating point** and it is chosen such the values of precision and recall




TP True Positives  
FP False Positives  
TN True Negatives  
FN False Negatives

# Criteria for model training

- Evaluation metrics for classification tasks

- **Precision/recall tradeoff:**

One typically has to choose between high precision or high recall – increasing one reduces the other



Source: A. Geron, Hands-On Machine Learning with Scikit-Learn, Keras & Tensorflow

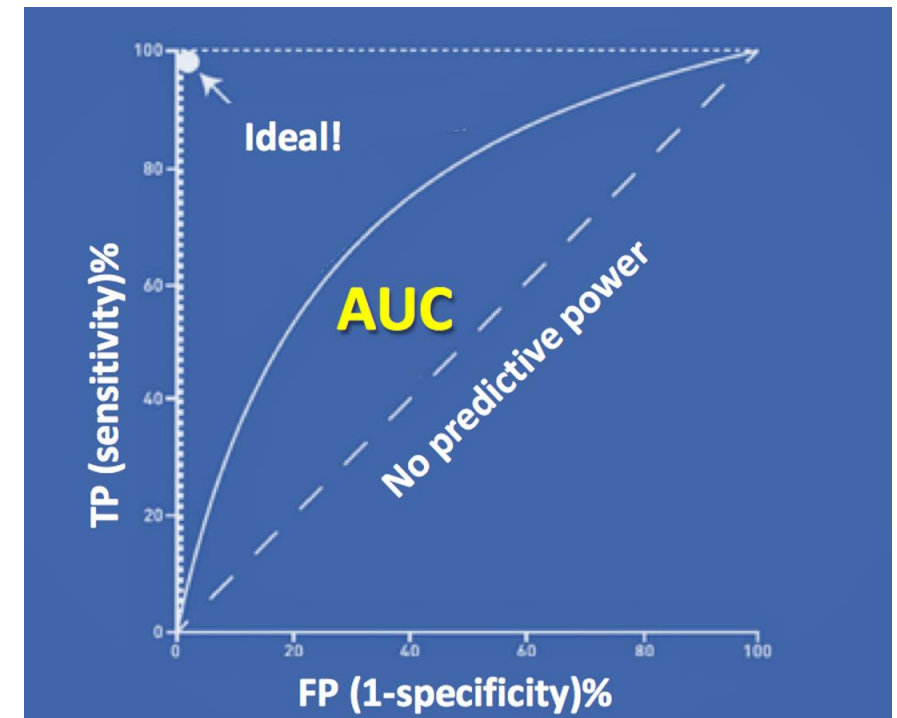
# Criteria for model training

- Evaluation metrics for classification tasks
- **receiver operating characteristic (ROC) curve**: a very useful tool for binary classification

- Y-axis: **true-positive rate (TPR)** – also known as recall
- X-axis: **false-positive rate (FPR)**

$$TPR = \frac{TP}{TP + FN} = \text{recall} \qquad FPR = \frac{FP}{FP + TN}$$

- **Area under the curve (AUC)**:
- a threshold-independent metric for binary classification!



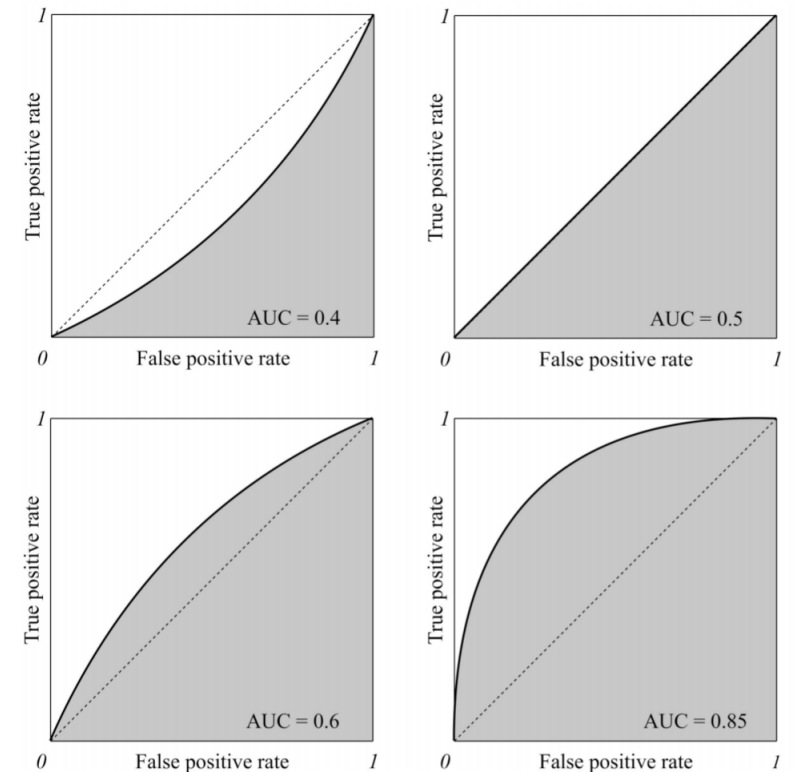
# Criteria for model training

- Evaluation metrics for classification tasks
- **receiver operating characteristic (ROC) curve**: a very useful tool for binary classification

- Y-axis: **true-positive rate (TPR)** – also known as recall
- X-axis: **false-positive rate (FPR)**

$$TPR = \frac{TP}{TP + FN} = \text{recall} \quad FPR = \frac{FP}{FP + TN}$$

- **Area under the curve (AUC)**:
- a threshold-independent metric for binary classification!



Source: A. Burkov, Machine Learning Engineering

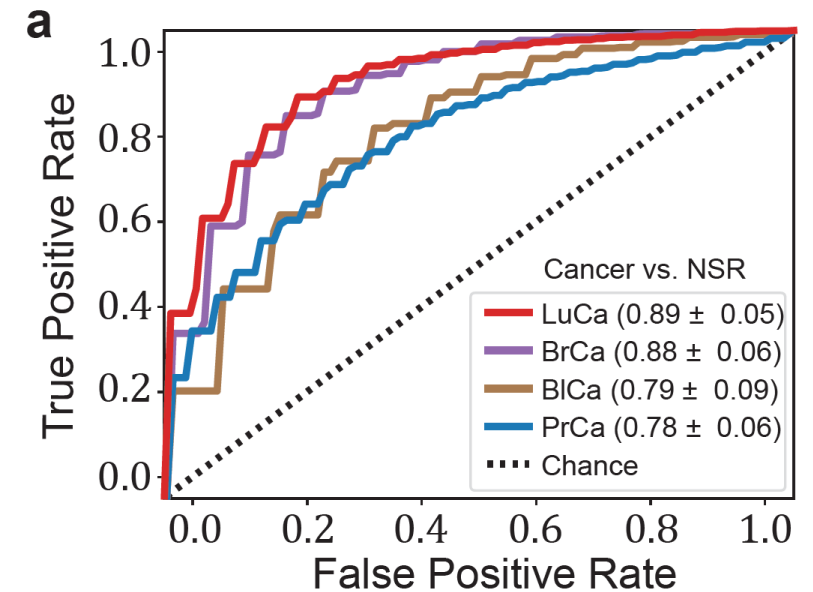
# Criteria for model training

- Evaluation metrics for classification tasks
- **receiver operating characteristic (ROC) curve**: a very useful tool for binary classification

- Y-axis: **true-positive rate (TPR)** – also known as recall
- X-axis: **false-positive rate (FPR)**

$$TPR = \frac{TP}{TP + FN} = \text{recall} \quad FPR = \frac{FP}{FP + TN}$$

- **Area under the curve (AUC)**:
- a threshold-independent metric for binary classification!



Train, validation and test data sets




# Train, validation and test data sets

- Only a subset of the available data can be used for training – the ability of a trained network to generalize must be evaluated on independent data
  1. *Train set:*
    - Used for model training
  2. *Validation set:*
    - Used for evaluating prediction quality
    - Improvements are made to the model afterwards
    - The network is retrained and reevaluated in an iterative procedure
  3. *Test set:*
    - Used as a final test of the prediction performance
    - It cannot influence any decisions about the model

# Train, validation and test data sets

- **cross-validation** (CV): a more advance method for evaluating network's performance
  - Split data into  $k$  equally-sized partitions
  - Use each part as a test set and combine the  $k - 1$  others for training
  - Obtain  $k$  results and average them

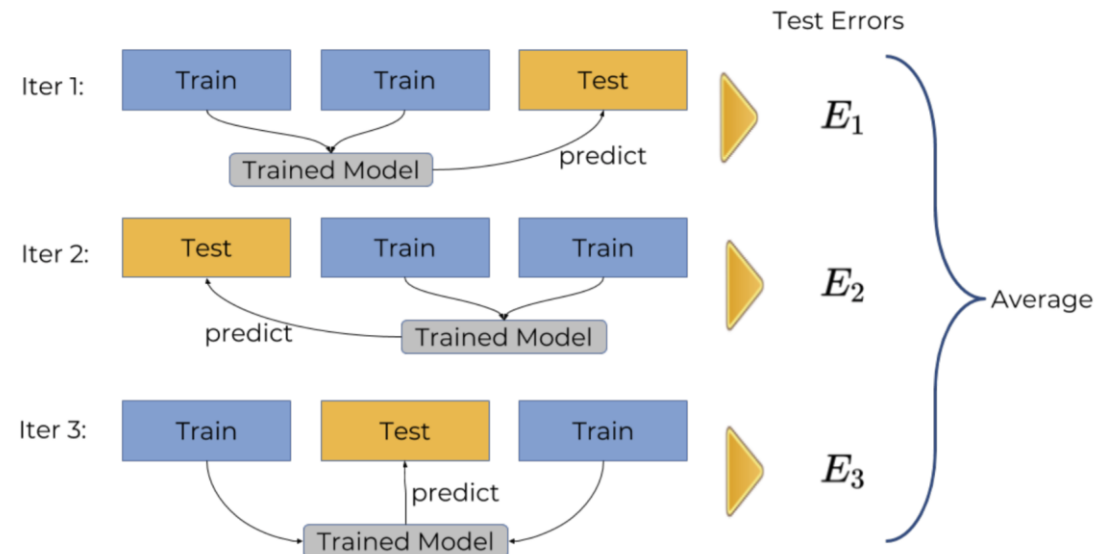


# Train, validation and test data sets

- **cross-validation** (CV): a more advance method for evaluating network's performance

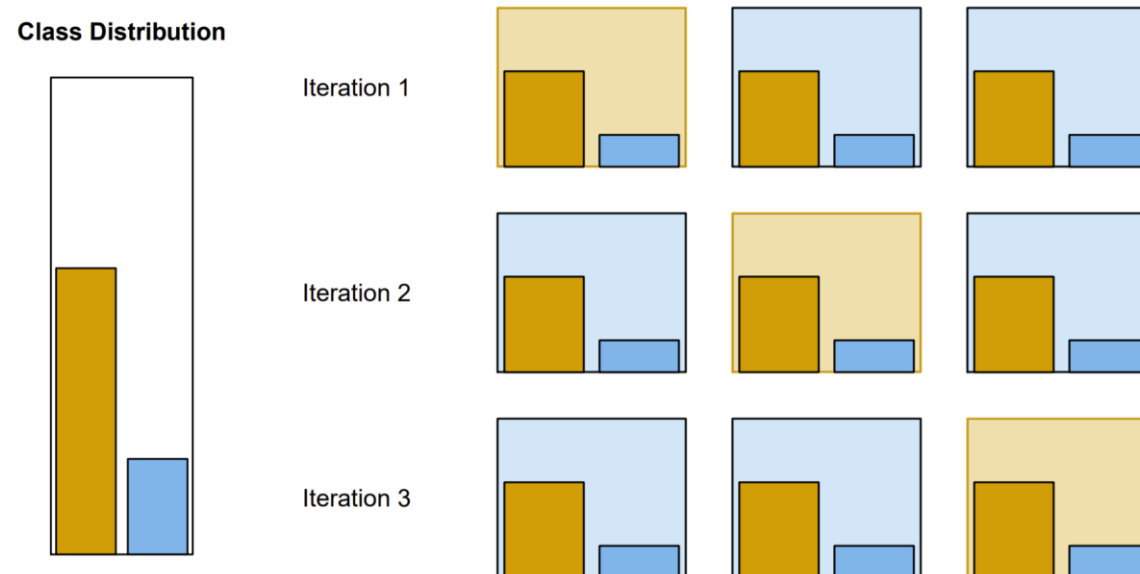
- Split data into  $k$  equally-sized partitions
- Use each part as a test set and combine the  $k - 1$  others for training
- Obtain  $k$  results and average them

- Example: 3-fold cross-validation



# Train, validation and test data sets

- **cross-validation** (CV): a more advance method for evaluating network's performance
  - K-fold cross-validation can be performed in a stratified way




# Train, validation and test data sets

- **cross-validation** (CV): a more advance method for evaluating network's performance
  - $k = 5$  or  $10$  are common choices: they use 80% or 90% of the data for training
  - For  $k = N$  we obtain a method known as **leave-one-out** (LOO)
    - $N-1$  data points for training and 1 for testing
  - Performance estimates tend to be pessimistically biased (as the size of the training sets is smaller than  $n$  and we learn less)
  - This bias increases as  $k$  gets smaller. LOO is nearly unbiased, but has high variance
  - Repeated  $k$ -fold CV (multiple random partitions) can improve error estimation for small sample size.

# Monitoring


# Monitoring

- Objective functions during training can provide useful information about the quality of model – can be plotted as function of epochs
  - Evaluate objective function on both the *train* and *validation* sets
  - Gap between them is referred to as **generalization error**



# Monitoring

- Objective functions during training can provide useful information about the quality of model – can be plotted as function of epochs
  - Evaluate objective function on both the *train* and *validation* sets
  - Gap between them is referred to as **generalization error**
  - For classification tasks: *accuracy* or *AUC* are used instead





# Regularization

# Regularization

- It is possible that optimization gets stuck at an “unwanted” *local minimum*, resulting into overfitting – several **regularization** methods exist for overcoming this issue
  - Problem with “unwanted” local minima is that the network could learn properties in the training data that are not generally valid
  - Definition: “*Regularization is any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error*” I. Goodfellow et al.

# Regularization

- Regularization methods:

## 1. **Early stopping**

- Critical point in training when/if the value of the objective function evaluated on the validation set starts rising – model starts to overfit
- At this point, training is terminated to avoid further overfitting

# Regularization

- Regularization methods:

## 2. Norm penalties

- Add a “penalty” term in the objective function to prevent the formation of large-valued parameter weights  $W$  – penalties are parameter-dependent
- Two variations:

1.  $L_1$  norm  $\mathcal{L} = \mathcal{L}_{MSE} + \sum_{i=1}^N \sum_{j=1}^N |W_{i,j}|$

1.  $L_2$  norm  $\mathcal{L} = \mathcal{L}_{MSE} + \sum_{i=1}^N \sum_{j=1}^N W_{i,j}^2$


# Regularization

- Regularization methods:

## 2. Norm penalties

1.  $L_1$  norm

$$\mathcal{L} = \mathcal{L}_{MSE} + \sum_{i=1}^N \sum_{j=1}^N |W_{i,j}|$$




# Regularization

- Regularization methods:

## 2. Norm penalties

1.  $L_2$  norm

$$\mathcal{L} = \mathcal{L}_{MSE} + \sum_{i=1}^N \sum_{j=1}^N W_{i,j}^2$$



# Regularization



- Regularization methods:

## 3. Dropout

- Individual nodes are temporarily-dropped at random during training
- typical dropout rates: 0.2, ..., 0.5
- It forces the network to create different mappings
- This method improves the network's stability
- For predictions on new data, all nodes are used

- Example situation:

- Sometimes network layers co-adapt to correct mistakes from prior layers
- This can potentially be broken up with dropout



# Regularization

- Regularization methods:

## 4. Additional regularization methods

- Collect more data!
- **Data augmentation:**
  - Generate additional data by modifying the already existing. E.g. rotated images
- **Noise injection:**
  - Alter input values during training using random noise – leads to stabilization
- **Ensemble training:**
  - Train many different networks on the same data for the same problem and combine their predictions – increases reliability



# Hyperparameters

# Hyperparameters


- List of **hyperparameters** to be tuned and decisions to be made
  - Number of hidden layers
  - Numbers of nodes per layer
  - Activation functions
  - Network initialization
  - Type and strength of regularization
  - Size of the minibatches
  - Learning rate and learning strategy

# Hyperparameters

- Methods for **hyperparameter tuning**

## 1. Grid search

- Create a grid where each point corresponds to particular combinations of hyperparameter values
- Evaluate model performance for each point using the objective function values on the validation set
- Good method if the hyperparameter space is not too large




# Hyperparameters

- Methods for **hyperparameter tuning**

## 2. Random search

- Instead of a grid, provide statistical distributions for each hyperparameter from which values are randomly sampled
- More efficient than grid search
- Works really well!




# Hyperparameters

- Methods for **hyperparameter tuning**

## 3. Coarse-to-fine search

- Random search followed by grid search
- After finding the high-potential regions, investigate the area around it and fine tune the values
- How many regions to investigate: depends on the time available



# Hyperparameters

- Methods for **hyperparameter tuning**



- Other methods

- **Bayesian** techniques: Differ from grid or random searches in that they use information from past evaluations to choose the next hyperparameter values to evaluate. Bayesian methods can learn
    - **Gradient-based** techniques...
    - **Evolutionary-optimization** techniques...
    - Other algorithmic methods...

# Activation functions

# Activation functions

- One of the most important decisions concerns the choice of the activation function
  - Good first choice: **ReLU** activation
  - **Leaky ReLU**: extension of ReLU with negative results also passed on
  - **Sigmoid** and **hyperbolic tangent (tanh)** function are rarely used in hidden layers because they can lead to vanishing gradient (training stops). They are used in the final output layer to constrain the result into  $[0,1]$  or  $[-1,1]$  respectively
- More detailed discussion on activation functions will be presented in later lectures





# Activation functions

a) ReLU


b) Leaky ReLU

c) Sigmoid


d) Tanh

e) ELU


f) SELU




(a)




(b)




(c)



(d)



(e)



(f)

# Summary

- The entire data set is usually split into three sets: **train**, **validation** and **test** sets
- A network's ability to **generalize** means that it can provide correct predictions on data that were not used for training/optimizing it
- The **bias-variance tradeoff** shows that reducing bias on model's predictions comes with the expense of increasing its variance and vice versa
- **Overfitting** is what happens when a network gives good predictions on the train set but fails to generalize, while **underfitting** is when the network fails to capture the complexity of the underlying distribution that generates the data
- The training procedure is typically **monitored** using the behavior of the objective function as function of the number of epochs
- **Regularization** methods help reduce overfitting and stabilize network training
- **Hyperparameters** define the type of a network and how it is trained